

SET REINFORCEMENT LEARNING: PRINCIPLES AND APPROACHES

Jubayer Ibn Hamid, Ifdita Hasan Orney, Chelsea Finn, Dorsa Sadigh and Noah Goodman

Stanford University

Correspondence to: jubayer@stanford.edu

ABSTRACT

In this note, we discuss the set reinforcement learning framework, various objectives that can be optimized in this setting, and develop apparatus for theoretical analysis. The overarching focus of this note is to understand the ways in which set reinforcement learning couples generations within a set and couples multiple sub-objectives within the main learning objective. Towards that end, we work on developing a set of definitions that characterize this coupling for arbitrary objectives.

CONTENTS

1	Introduction	2
2	Set Reinforcement Learning Framework	2
2.1	Sequential Formulation	2
2.2	Single-turn Formulation	4
2.3	General Algorithm	5
2.4	Polychromic Objective	8
2.5	Discussion on leave-one-out advantages	10
3	Scaffold and Marginal Advantage	17
3.1	Scaffold Value	17
3.2	Marginal Set Advantage	18
3.3	Entropy Analysis	19
3.4	Analysis of Polychromic Objectives	20
4	Ascendant Value and Oversampling	23
5	Stratification	28
6	Application: Exploration	32
6.1	Poly-PPO	33
6.2	Poly-EPO	34
6.3	Spiral	34

1. INTRODUCTION

Reinforcement learning (RL) provides a general framework for training agents through interaction with an environment enabling them to learn from their own experiences. Its flexibility stems both from the diversity of algorithmic approaches and from the ability to specify the environment and reward function. In standard formulations, the reward is defined over individual actions or trajectories, and the goal is to learn a policy that selects an optimal action at each state. However, in many settings, the objective is not to produce a single optimal action, but rather a *set* of actions or trajectories with desirable collective properties. For example, in language model (LM) post-training, one often samples multiple candidate generations before applying verification, ranking, search procedures or self-aggregation. In such cases, performance depends on properties of the entire set rather than on individual generations.

Set reinforcement learning (set RL), introduced in [HOX⁺26], generalizes the RL framework by optimizing policies with respect to objectives defined over sets of trajectories. Concretely, given a policy that induces a distribution over sets of trajectories, set RL aims to maximize a set-level objective function that couples elements within the set under a shared learning signal. This perspective provides a high degree of flexibility in designing objectives, analogous to reward design in standard RL, while enabling fundamentally different forms of learning signals.

A central feature of set RL is the *coupling* between elements of a set. Unlike standard RL, where each trajectory contributes independently to the objective, set-level objectives can induce dependencies across trajectories, leading to learning signals that reflect collective properties. Understanding how this coupling manifests is a primary focus of this note.

In this note, we investigate the structure and implications of set RL objectives. We analyse several classes of objectives, develop tools for reasoning about their learning signals, and highlight key phenomena arising from coupling within sets. Our discussion is motivated primarily by applications in LM post-training, where set-based objectives naturally arise in the context of sampling, search, and test-time compute.

2. SET REINFORCEMENT LEARNING FRAMEWORK

We start by recalling the standard reinforcement learning framework. Consider a Markov decision process defined by state space $\mathcal{S}$, action space $\mathcal{A}$, transition dynamics distribution $p(s_{t+1} \mid s_t, a_t)$, reward function $r : \mathcal{S} \times \mathcal{A} \rightarrow \mathbb{R}$, initial state distribution ρ_0 and discount factor $\gamma \in (0, 1)$. In reinforcement learning, we aim to find the policy π_θ parametrized by $\theta \in \Theta$ that solves the following problem:

$$\max_{\theta} \mathbb{E}_{s_0 \sim \rho} \mathbb{E}_{\tau \sim \pi_\theta} \left[\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t) \mid s_0 \right]. \quad (1)$$

Here, a trajectory τ refers to a tuple $(s_0, a_0, r_0, s_1, a_1, r_1, \dots)$. Note that the notation here is overloaded; we write $\tau \sim \pi_\theta(\cdot \mid s_0)$ to mean that each action $a_t \sim \pi_\theta(\cdot \mid s_t)$ is sampled from the policy while the states are sampled from the MDP’s transition dynamics distribution, i.e. $s_{t+1} \sim P(\cdot \mid s_t, a_t)$. Intuitively, the problem here is to find the policy that maximizes the expected returns from a trajectory τ sampled under the policy. One can think of this problem as essentially aiming to solve $\max_{\theta} \mathbb{E}_{s_0 \sim \rho, \tau \sim \pi_\theta(\cdot \mid s_0)} [f(\tau)]$ where $f(\tau)$ is the (discounted) sum of rewards in the trajectory τ . We will refer to this framework as *standard reinforcement learning* for clarity.

Set reinforcement learning is a framework that, at a high level, attempts to find that policy that maximizes the expected value of a *set* of trajectories sampled under the policy with respect to some objective function. We first look at the idealised sequential formulation setting and then study the more practical setting.

2.1. SEQUENTIAL FORMULATION

In the idealised sequential formulation, at each visited state s , the policy samples a set of n actions independently from $\pi_\theta(\cdot \mid s)$ where $n > 1$ is a fixed integer. Given $a_{1:n} := (a_1, \dots, a_n) \stackrel{\text{i.i.d.}}{\sim} \pi_\theta(\cdot \mid s)$, a set-level reward $f(s, a_{1:n})$ is assigned to the entire set. Repeating this procedure recursively at each timestep

generates an n -ary rollout tree induced by our policy. At depth t , the tree contains n^t states, which we denote by $s_t^{(1)}, \dots, s_t^{(n^t)}$. For each such state $s_t^{(i)}$, let $(a_t)_{1:n}^{(i)}$ denote the n -tuple of actions sampled from $\pi_\theta(\cdot | s_t^{(i)})$. The goal is to solve the following optimization problem defined using set value function:

$$\max_{\theta} \mathbb{E}_{s_0 \sim \rho} \mathbb{E}_{\pi_\theta} \left[\sum_{t=0}^{\infty} \sum_{i=1}^{n^t} \gamma^t f(s_t^{(i)}, (a_t)_{1:n}^{(i)}) | s_0 \right]. \quad (2)$$

To be clear, at each state $s_t^{(i)}$, the set of n actions are sampled independently, i.e. $a_{t,1}^{(i)}, \dots, a_{t,n}^{(i)} \stackrel{i.i.d.}{\sim} \pi_\theta(\cdot | s_t^{(i)})$. In particular, we require that f is an objective such that $\mathbb{E}_{a_{1:n} \sim \pi_\theta(\cdot | s)} [f(s, a_{1:n})]$ cannot be written as $\mathbb{E}_{a \sim \pi_\theta(\cdot | s)} [g(s, a)]$ where g is independent of θ for all policy parameters $\theta \in \Theta$. In other words, we cannot collapse the set RL problem to a standard RL one. This is an important feature in the definition of set RL.

This caveat is particularly important. For example, if $f(s, a_{1:n}) = \frac{1}{n} \sum_{i=1}^n r(s, a_i)$, then the set RL problem in eq. (2) becomes equivalent to the standard RL problem in eq. (1). However, one can consider several objectives that satisfy the requirements for set RL; for example, $f(s, a_{1:n}) = \max_i r(s, a_i)$, which is the pass@ n objective.

Note that we will almost always use the notation $a_{1:n}$ to mean that it is the following multiset: $\{a_1, \dots, a_n\}$. The multiset is allowed to have multiple instances of the same element and the number of instances of any element is called its multiplicity.

To understand what set RL is attempting to optimize, it is useful to construct a notion of value functions in this framework.

Definition 2.1 (Set Value Functions). *Given a policy π generating a state-visitation tree and a set objective $f : \mathcal{S} \times \mathcal{A}^n \rightarrow \mathbb{R}$, the set state value function $V_\pi^\#(s; f)$ and the set state-action value function $Q_\pi^\#(s, a_{1:n}; f)$ are defined as*

$$V_\pi^\#(s; f) = \mathbb{E}_\pi \left[\sum_{t=0}^{\infty} \sum_{i=1}^{n^t} \gamma^t f(s_t^{(i)}, (a_t)_{1:n}^{(i)}) | s_0 = s \right], \quad (3)$$

$$Q_\pi^\#(s, a_{1:n}; f) = \mathbb{E}_\pi \left[\sum_{t=0}^{\infty} \sum_{i=1}^{n^t} \gamma^t f(s_t^{(i)}, (a_t)_{1:n}^{(i)}) |_{(a_0)_{1:n} = a_{1:n}}^{s_0 = s} \right] \quad (4)$$

Here, we assume that the discount factor $\gamma \in (0, \frac{1}{n})$ to ensure values remain bounded - the range is smaller than in standard RL since we add the expected sum of rewards from n actions stemming out of each state. Intuitively, $V_\pi^\#(s)$ is the expected discounted return of the entire state tree rooted at s , as illustrated in fig. 1, where the reward at each node is given by the set objective. In contrast, $Q_\pi^\#(s, a_{1:n})$ evaluates the expected return of the tree that begins at s with the specific action set $a_{1:n}$.

With these definitions, we first see that the set RL problem can be written as:

$$\max_{\theta} \mathbb{E}_{s_0 \sim \rho} [V_{\pi_\theta}^\#(s_0; f)].$$

In other words, in set reinforcement learning, we want to learn the policy that maximizes the expected return from a *tree* generated using our policy is maximized. In contrast, in standard reinforcement learning, we want to learn the policy such that the expected return from a *trajectory* is maximized.

These constructions allow us to translate several well-known ideas and algorithms from standard reinforcement learning to the set reinforcement learning framework. As an example, we have the following result, which is the set RL analogue of the performance-difference lemma from [KL02b]:

Lemma 2.2. *Given any two policies π_θ and π_β and a fixed initial state s_0 , under any set objective function f ,*

$$V_{\pi_\theta}^\#(s_0; f) - V_{\pi_\beta}^\#(s_0; f) = \frac{1}{1 - \gamma n} \mathbb{E}_{s \sim d_{\pi_\theta}^\#(\cdot), a_{1:n} \sim \pi_\theta(\cdot | s)} [A_{\pi_\beta}^\#(s, a_{1:n}; f)].$$




Figure 1: The set value of a state (circled) is the expected discounted return of the subtree (highlighted) rooted in this state.

Here $d_{\pi}^{\sharp}(s)$ is the stationary state-visitation distribution induced by π_{θ} . Similar to the standard reinforcement learning setting, this result says that if we update the policy π_{β} to π_{θ} such that, at all states visited by π_{θ} , the advantage $A_{\pi_{\beta}}^{\sharp}(s, a_{1:n}; f) = Q_{\pi_{\beta}}^{\sharp}(s, a_{1:n}; f) - V_{\pi_{\beta}}^{\sharp}(s)$ of a set of actions taken by our new policy π_{θ} is positive, then we will get a policy that has strictly higher performance (as measured by its value). This suggests that many of the principles underlying methods, like PPO from [SWD⁺17], can be extended to the set reinforcement learning paradigm as well for policy improvement. Similarly, algorithms such as Q-learning can also be adapted provided that we sample rollouts in the form of state-visitation trees.

The sequential formulation is conceptually useful, but it is computationally impractical in most settings because the number of sampled states grows exponentially with depth. However, sampling complexity aside, this formulation does indeed capture the goal of several problems. As an example, when training an agent to play a game such as chess, we want the model to sample the optimal *tree* of states visited where the tree can be scored by the sum of rewards received at its leaf nodes.

Throughout the rest of this paper, we will concern ourselves with significantly more practical formulations, such as the single-turn formulation in the next section. However, it is important to remind ourselves that all such formulations are approximations of the core ambition of the set reinforcement learning framework: to train a model to sample the optimal tree of trajectories under some set-level objective function.

2.2. SINGLE-TURN FORMULATION

In many practical settings, we do not want to sample state-visitation trees like we did in the sequential framework because of the large sample complexity and other practical challenges. For example, in the language model (LM) post-training setting, we sample a batch of prompts $x_1, \dots, x_{|B|} \sim \mathcal{D}$ where $\mathcal{D}$ is our training set, and then sample N generations from the LM conditioned on each prompt, i.e. $y_{1:n} := (y_1, \dots, y_N) \sim \pi_{\theta}(\cdot | x_i)$ for each x_i .

In this setup, set RL aims to solve:

$$\max_{\theta} \mathbb{E}_{x \sim \mathcal{D}} \mathbb{E}_{y_{1:n} \sim \pi_{\theta}(\cdot | x)} [f(x, y_{1:n})]. \quad (5)$$

Again, we require that the set objective function f is such that we cannot write, for each x , $\mathbb{E}_{y_{1:n}} [f(x, y_{1:n})]$ as $\mathbb{E}_y [g(x, y)]$ where g is independent of θ for all policy parameters $\theta \in \Theta$. This requirement is crucial because, otherwise, under some objectives (for e.g., $f(x, y_1, \dots, y_n) = \frac{1}{n} \sum_{i=1}^n r(x, y_i)$) the set RL formulation in eq. (5) becomes equivalent to a standard RL formulation using some non-differentiable reward function.

We now look at some examples of objectives that satisfy the set reinforcement learning framework.

Example 2.3. Let $f(x, y_{1:n}) = \max_i r(x, y_i)$. This objective corresponds to pass@ n optimization [WK25].

Example 2.4. More generally, one can consider objectives that represent the usage of test-time compute. Let $\text{Alg}(x, y_1, \dots, y_n) \in \mathcal{Y}$, i.e. the algorithm returns the final generation to return. Then, one can consider $f(x, y_{1:n}) = r(x, \text{Alg}(x, y_{1:n}))$ which is an objective that rewards the generation chosen by our test-time compute. Alternatively, if the algorithm returns a subset of the generations, i.e. $\text{Alg}(x, y_{1:n}) \in \mathcal{Y}^m$ for $m < n$, then we can set $f(x, y_{1:n}) = \mu(\{r(x, \text{Alg}(x, y_{1:n}))_j\}_{j=1}^m)$ which is the average reward in the set of generations selected by the algorithm.

Example 2.5. Let $d(x, y_{1:n}) \in [0, 1]$ be a measure of the diversity in the generations $(y_1, \dots, y_n)$. Then, the following objective also satisfies set RL requirements: $f(x, y_{1:n}) = \frac{1}{n} \sum_{i=1}^n r(x, y_i) \cdot d(x, y_{1:n})$. This is the *polychromic objective*, designed to encourage the model to balance exploration and exploitation, which we will look at in detail later.

To consolidate our understanding of how set RL optimizes differently from standard RL, we look at the policy gradient. After applying the score-function identity to the joint density, we get the following policy gradient:

$$\begin{aligned} & \nabla_{\theta} \mathbb{E}_{x \sim \mathcal{D}} \mathbb{E}_{y_{1:n} \sim \pi_{\theta}(\cdot | x)} [f(x, y_{1:n})] \\ &= \mathbb{E}_{x \sim \mathcal{D}} \mathbb{E}_{y_{1:n} \sim \pi_{\theta}(\cdot | x)} \left[(f(x, y_{1:n}) - \hat{f}(x)) \sum_{i=1}^n \nabla_{\theta} \log \pi_{\theta}(y_i | x) \right] \end{aligned} \quad (6)$$

where $\hat{f}(x)$ is a set-level baseline. In other words, $\hat{f}(x)$ is independent of $y_{1:n}$ or any strict subset of $y_{1:n}$. The defining feature of (6) is that all generations in the sampled set $y_{1:n}$ share the same scalar learning signal, $f(x, y_{1:n}) - \hat{f}(x)$. In contrast, in the standard RL setting, we get:

$$\nabla_{\theta} \mathbb{E}_{x \sim \mathcal{D}} \mathbb{E}_{y \sim \pi_{\theta}(\cdot | x)} [r(x, y)] = \mathbb{E}_{x \sim \mathcal{D}} \mathbb{E}_{y \sim \pi_{\theta}(\cdot | x)} [\nabla_{\theta} \log \pi_{\theta}(y | x) \cdot (r(x, y) - b)] . \quad (7)$$

As such, in standard RL, each generation receives its own individual advantage. In set RL, the objective f couples the samples in the set, and the gradient assigns equal credit to all elements in the set through a shared set-level signal. For this reason, admissible baselines must be independent of the sampled set $y_{1:n}$ or any subset of it.

The notion of coupling will be a recurring theme in this note. So far we have seen that in set RL, the objectives within a set are coupled in the sense that they all receive the same score. We will soon see that, depending on the choice of the function f , multiple *objectives* can be “coupled”, i.e. sets can be optimized to balance several goals for our learning agent.

A naive algorithm for training an agent in this set RL framework would be to sample multiple sets of size n per prompt $x \sim \mathcal{D}$ to get a Monte Carlo estimate of the gradient in eq. (6). However, the sample complexity of this naive algorithm could become quite large. In particular, if the baseline, $\hat{f}(x)$ is also a Monte Carlo estimate, one would need to sample several sets of size n for there to be a meaningful reduction in variance. We will now look at a more practical, general recipe for training a policy using set reinforcement learning.

2.3. GENERAL ALGORITHM

We now discuss a general recipe for optimizing a model using set reinforcement learning, as introduced in [OHR⁺26]. This will be strongly motivated by the language model (LM) post-training setting.

Suppose, that our set objective, $f(x, y_1, \dots, y_n)$, is defined over sets of size n . Moreover, we assume that f is symmetric in the arguments $y_{1:n}$. If we wanted to use an empirical estimate (e.g., a Monte Carlo estimate) of the gradient in (6), then we would need to sample several sets of size n conditioned on a single

prompt which can be computationally expensive. In this section, we consider how we can instantiate set reinforcement learning using a computationally feasible gradient estimator.

For a fixed prompt x , we first draw N i.i.d. generations $y_1, \dots, y_N \sim \pi_\theta(\cdot | x)$. We require that the number of samples, N , is greater than the size of sets, n , in the set RL framework. Next, from these samples, we construct K sets, $G_1, \dots, G_K$, of size $n < N$, in a combinatorial fashion without replacement:

$$G_j = \{y_{(j_1)}, \dots, y_{(j_n)}\}, \quad j_1 < \dots < j_n, \{j_1, \dots, j_n\} \subseteq \{1, \dots, N\}.$$

The maximum number of sets that can be constructed in this manner is $K = \binom{N}{n}$. Using all $\binom{N}{n}$ sets can reduce the variance of the resulting gradient estimator. When evaluating f is computationally expensive, scoring all $\binom{N}{n}$ sets may be impractical. In this case, we instead sample K sets uniformly (without replacement) from the collection of all $\binom{N}{n}$ possible sets. Either of these choices leads us to an unbiased estimator of the set RL gradient as we will see later.

Now that we have constructed K sets, for each constructed set G_j , we compute the set-level score under the set objective function:

$$f(x, G_j) = f(x, y_{(j_1)}, \dots, y_{(j_n)}).$$

We construct a variance-reduction baseline using the sampled sets:

$$\hat{f}(x) := \frac{1}{K} \sum_{j=1}^K f(x, G_j).$$

In other words, the baseline is simply an average of the scores of all sets under our objective function f . Given the baseline, for each set G_j , we can now construct a set advantage function which is simply

$$\widehat{A^\#}(x, G_j; f) = f(x, G_j) - \hat{f}(x). \quad (8)$$

So far, we have computed everything at the set-level, i.e. we constructed sets, scored them and computed each set's advantage. We now derive an advantage function at the level of a single generation.

Marginal set advantage of a single generation. Intuitively, we will set the advantage of each generation to be simply the sum of the advantages of all sets containing the generation. Let $\mathcal{G}(y)$ be the collection of all sets that contain the fixed generation y . Note that $\mathcal{G}(y) = \{G \in G_{1:K} \mid y \in G\} \subset G_{1:K}$. We define the marginal set advantage of the generation y to be:

$$\widehat{A_{\text{marg}}^\#}(x, y; f) := \frac{1}{|\mathcal{G}(y)|} \sum_{G \in \mathcal{G}(y)} A^\#(x, G; f). \quad (9)$$

Final Estimator. Using this, we use the following set RL gradient estimator:

$$\nabla_\theta \widehat{\mathbb{E}}_{x \sim \mathcal{D}, y_{1:n} \sim \pi_\theta(\cdot | x)} [f(x, y_{1:n})] = \widehat{\mathbb{E}}_{x \sim \mathcal{D}, y_1, \dots, y_N \sim \pi_\theta(\cdot | x)} \left[\sum_{i=1}^N \nabla_\theta \log \pi_\theta(y_i | x) \widehat{A_{\text{marg}}^\#}(x, y_i; f) \right] \quad (10)$$

Now, we have all the ingredients required for optimizing an LM with respect to the set RL framework. Since the marginal set advantage gives us an advantage function over single generations, we can use existing *standard* RL algorithms, such as PPO [SWD⁺17] or GRPO [SWZ⁺24], to optimize the set objective by replacing their standard advantage, $A(x, y_i)$, with $\widehat{A_{\text{marg}}^\#}(x, y_i)$. The pseudocode for our general recipe is provided in algorithm 1.

Note that our estimator is an unbiased estimator of the true policy gradient of set RL (up to a scaling factor). This is true for both the case where we construct all $\binom{N}{n}$ sets and the case where we sample K sets uniformly without replacement. We show this in the following proposition:

Algorithm 1 Gradient of On-Policy Implementation of Set RL

Require: Set objective f , Policy π_θ , batch of inputs B , number of rollouts N , set size n

```

1: for each input  $x \in B$  do
2:   Sample  $y_1, \dots, y_N \sim \pi_\theta(\cdot | x)$ 
3:   // Construct sets either by enumerating all  $\binom{N}{n}$  subsets or by uniformly sampling  $K$  subsets from them
4:   Construct sets  $G_1, \dots, G_K$  from  $\{y_{1:N}\}$  without replacement
5:   Compute score of each set  $\{f(x, G_l)\}_{l=1}^K$  and baseline  $b = \frac{1}{K} \sum_{l=1}^K f(x, G_l)$ 
6:   Compute set advantage of each set  $\widehat{A}^\sharp(x, G_l; f) = f(x, G_l) - b$  for  $l = 1, \dots, K$ 
7:   // Let  $\mathcal{G}(y)$  be all the sets in  $G_1, \dots, G_K$  that contain  $y$ 
8:   Compute marginal set advantage  $\widehat{A}_{\text{marg}}^\sharp(x, y_j; f) := \frac{1}{|\mathcal{G}(y_j)|} \sum_{G \in \mathcal{G}(y_j)} \widehat{A}^\sharp(x, G)$  for  $j = 1, \dots, N$ 
9:    $\hat{g}(x) \leftarrow \sum_{j=1}^N \nabla_\theta \log \pi_\theta(y_j | x) \cdot \widehat{A}_{\text{marg}}^\sharp(x, y_j; f)$ 
10: end for
11:  $\hat{g} \leftarrow \frac{1}{|B|} \sum_{x \in B} \hat{g}(x)$ 
12: return  $\hat{g}$ 

```

Proposition 2.6. Fix a prompt x , and let $y_1, \dots, y_N \stackrel{\text{i.i.d.}}{\sim} \pi_\theta(\cdot | x)$ be our independently sampled N generations and let $f : \mathcal{X} \times \mathcal{Y}^{\oplus n} \rightarrow \mathbb{R}$ be our set objective. Then,

$$\mathbb{E} \left[\sum_{i=1}^N \nabla_\theta \log \pi_\theta(y_i | x) \widehat{A}_{\text{marg}}^\sharp(x, y_i; f) \right] = M \nabla_\theta \mathbb{E}_{y_{1:n} \sim \pi_\theta(\cdot | x)} N[f(x, y_{1:n})],$$

where $M \in \mathbb{R}_{>0}$ is a constant scaling factor that depends on the number of sets we construct. In particular, when we construct all $\binom{N}{n}$ sets without replacement, we have that

$$M = \binom{N}{n} - \binom{N-1}{n-1}$$

and when we sample, uniformly, K sets without replacement from $K_{\text{all}} := \binom{N}{n}$ sets, we have that

$$M = \frac{N}{n} q_K - 1,$$

where

$$q_K := \Pr(C_i(\mathcal{T}) > 0) = 1 - \frac{\binom{N-1}{n} \binom{K}{n}}{\binom{N}{n} \binom{K}{n}},$$

with the convention that $\binom{a}{K} = 0$ when $K > a$. Consequently, after also taking expectation over $x \sim \mathcal{D}$ and scaling the learning rate, the estimator is an unbiased estimator of the set RL gradient.

The proof for this can be found in [OHR⁺26]. Note that the proof shows that constructing all $\binom{N}{n}$ sets allows us to get the complete U-statistic estimator which, provably, has the lowest variance.

Next, we briefly consider another example where we address the following question: can we use the set RL framework where the set size itself is not fixed? Suppose we have an objective that can be defined on sets of variable size. How can we optimize this objective over sets of varying size?

Example 2.7. Suppose, we sample N generations in total and we seek to optimize some objective over sets of size $n_1, n_2, \dots, n_C \in \mathbb{Z}_{\geq 1}$. Here, we assume that the objective function f could be defined on any of these sets (with some abuse of notation). For example, the objective $f(x, y_{1:n}) = \max_i r(x, y_i)$ does not directly depend on the size of the set insofar as the set has at least one generation. We could generalize the ideas we have presented so far to address this problem setting. First, we construct sets of size n_1 and denote the full collection of sets of size n_1 by $\mathcal{G}(n_1)$. For example, we could compute all possible sets of size n_1

combinatorially from the N generations. Similarly, we get a collection of sets of size n_2 called $\mathcal{G}(n_2)$, and so on. Then, we can optimize using set RL by optimizing the following:

$$\mathbb{E}_{y_1, \dots, y_N \sim \pi_\theta(\cdot|x)} \left[\frac{1}{\binom{N}{n_1}} \sum_{Y \in \mathcal{G}(n_1)} f(x, Y) + \dots + \frac{1}{\binom{N}{n_C}} \sum_{Y \in \mathcal{G}(n_C)} f(x, Y) \right]$$

Finally, we consider the case where we might want to optimize a set objective function f that is not symmetric in its arguments. The good news is that our preceding development does not strictly require that f is symmetric. In particular, the set reinforcement learning gradient is still the same, except one must view the sampled set as an ordered tuple:

$$\nabla_\theta \mathbb{E}_{(y_{1:n}) \sim \pi_\theta(\cdot|x)} [f(x, y_1, \dots, y_n)] = \mathbb{E}_{(y_{1:n}) \sim \pi_\theta(\cdot|x)} \left[(f(x, y_{1:n}) - \hat{f}(x)) \sum_{i=1}^n \nabla_\theta \log \pi_\theta(y_i | x) \right].$$

As for the general recipe, we first need to note that the total number of sets of size n we can construct from N generations is $\frac{N!}{(N-n)!}$. In most practical cases, this will be significantly larger than the case for symmetric f which is $\binom{N}{n}$. As such, it is more feasible to rely on uniform sampling, without replacement, of K sets of size n from the $\frac{N!}{(N-n)!}$ sets. Note that any generation y can appear in at most $n \frac{(N-1)!}{(N-n)!}$. The marginal set advantage of each generation should still be calculated in the same manner as before while keeping in mind that f is not symmetric and, as such, two different permutations of the ordered tuple are not necessarily equal in value under f . In particular, if we construct all possible sets, the marginal set advantage is

$$A_{\text{marg}}^\sharp(x, y_i; f) = \frac{1}{n \frac{(N-1)!}{(N-n)!}} \sum_{p=1}^n \sum_{(i_1, \dots, i_n) \in \tau, i_p = i} (f(x, y_{i_1}, \dots, y_{i_n}) - \hat{f}(x))$$

where $\tau = \{(i_1, \dots, i_n) \mid \text{each } i_j \in \{1, \dots, N\}, i_j \neq i_k \text{ for any } j \neq k\}$. In the case where we uniformly sample, without replacement, K sets from the $\frac{N!}{(N-n)!}$ sets, the marginal set advantage is simply the average set advantage of all sets containing the generation. Furthermore, Proposition 2.6 still holds in this case, although the scaling factor M changes. In the uniformly sampling case, this scaling factor becomes $M = \frac{N}{n} q_K - 1$ where

$$q_K := \Pr(C_i(\mathcal{S}_K) > 0) = 1 - \frac{\binom{(N-1)_n}{K}}{\binom{(N)_n}{K}},$$

where

$$(N)_n := \frac{N!}{(N-n)!}.$$

2.4. POLYCHROMIC OBJECTIVE

$$\max_{\theta} \mathbb{E}_{y_{1:n} \sim \pi_\theta(\cdot|x)} \left[\mathbb{E}_{y_* \sim \pi_\theta(\cdot|x, y_{1:n})} [r(x, y_*)] \right].$$

Exploration is a critical component in learning from experience. Firstly, it is necessary if you want to train an agent in a complex problem setting where one must explore to discover at least one successful solution; note that if the agent fails to discover any behaviour that gets some reward signal, then policy gradient methods will simply not be able to optimize at all since the advantage of all actions will be 0. Secondly, often it is quite important for the agent to discover several, diverse successful behaviours. An agent that is learning to play chess should ideally master multiple opening lines rather than relying on a single strategy that may prove ineffective against a different opponent at test time. Thirdly, exploration is necessary for an agent to stress test what it has learned and its internal beliefs. Otherwise, the model overfit on spurious correlations between the reward signal and specific behaviours.

The exploration-exploitation tradeoff in reinforcement learning is a well-known problem. This problem becomes especially pronounced in policy gradient methods. One must train the agent to learn what gets high

rewards (i.e. exploit) while also attempting novel strategies (i.e. explore). This requires a form of *optimism under uncertainty*: the agent must attempt strategies during data collection that are exploratory even if they do not immediately get high rewards. Since the data is sampled from the policy itself, the learning signal itself must encourage allocating probability mass to strategies that have not yet demonstrated high reward. For example, when attempting to solve an algorithmic problem, an agent exploring a dynamic programming formulation should continue refining that strategy even if its initial implementations fail unit tests, rather than immediately reverting to a familiar but less scalable exhaustive search approach.

However, the learning signal cannot simply encourage attempting novel strategies. It must also encourage the policy to exploit by re-attempting what has already receive high rewards. As such, the algorithm must address the *trade-off* between exploration and exploitation. We seek an optimization objective for reinforcement learning that satisfies the following desiderata. First, the objective should encourage *optimistic exploration* by assigning positive learning signal to trajectories that attempt novel reasoning strategies, even when those strategies have not yet yielded high reward. Second, the objective should *intrinsically encode the exploration–exploitation trade-off* and incentivize balancing these two goals. This allows the model itself to learn how to balance these competing objectives through optimization.

Polychromic objectives are a family of objectives that are designed to do that. We first look at a practical example of a polychromic objective in the same framework as section 2.2:

$$f_{\text{poly}}(x, y_1, \dots, y_n) = \frac{1}{n} \sum_{i=1}^n r(x, y_i) \cdot d(x, y_1, \dots, y_n), \quad (11)$$

where $d(x, y_1, \dots, y_n)$ is a function that measures the diversity of the set of generations, $y_{1:n}$. Here, we assume that $r(x, y) \in [0, 1]$ and $d(x, y_{1:n}) \in [0, 1]$. Because the set-RL gradient uses a shared advantage for all trajectories in a set, this objective increases the likelihood of successful behaviours *and* diverse exploratory trajectories. The shared advantage term amplifies exploratory trajectories that do not (yet) yield high rewards, pushing the policy to discover diverse strategies.

Before we move on to analysis, let us consolidate our understanding of how set reinforcement learning couples the generations in a set under this objective. One way to do this would be to rewrite the objective in the set RL framework as follows:

$$\begin{aligned} & \mathbb{E}_{y_{1:n} \sim \pi_\theta(\cdot|x)} \left[\frac{1}{n} \sum_{i=1}^n r(x, y_i) d(x, y_{1:n}) \right] \\ &= \frac{1}{n} \sum_{i=1}^n \mathbb{E}_{y_{1:n} \sim \pi_\theta(\cdot|x)} [r(x, y_i) \cdot d(x, y_{1:n})] \\ &= \frac{1}{n} \sum_{i=1}^n \mathbb{E}_{y_i \sim \pi_\theta(\cdot|x)} [r(x, y_i) \mathbb{E}_{y_{1:i-1}, y_{i+1:n} \sim \pi_\theta(\cdot|x)} [d(x, y_{1:n}) \mid y_i]] \\ &= \mathbb{E}_{y \sim \pi_\theta(\cdot|x)} [r(x, y) \mathbb{E}_{y_{1:n-1} \sim \pi_\theta(\cdot|x)} [d(x, y, y_{1:n-1})]] \end{aligned}$$

Note that, even though the outer expectation is over a single generation, the final expression is not equivalent to $\mathbb{E}_{y \sim \pi_\theta(\cdot|x)} [g(x, y)]$ for any g that is independent of θ . Consequently, taking the gradient, with respect to θ would not only contribute a learning signal for y sampled in the outer expectation, but also for $y_{1:n-1}$ in the inner expectation. This is simply another way of saying that the set score acts as the same learning signal for all generations in the set.

In contrast, [LZY⁺25] proposed optimizing $\mathbb{E}_{y \sim \pi_\theta(\cdot|x)} [r(x, y) \cdot d(x, y \mid y_{1:n-1})]$, but $y_{1:n-1}$ are treated as constants which turns $d(x, y \mid y_{1:n-1})$ into a θ -independent reward function. Consequently, optimizing this does not lead to a set-level learning signal being shared by all constituents whereas set reinforcement learning causes the polychromic objective’s signal to be shared by all generations in a set.

2.5. DISCUSSION ON LEAVE-ONE-OUT ADVANTAGES

In this section, we briefly discuss an alternative route prior works have taken to optimize set-level rewards. As noted by [TZSM25], one can use a leave-one-out advantage estimator as a replacement for the set RL gradient:

$$\begin{aligned} & \nabla_{\theta} \mathbb{E}_{y_{1:n} \sim \pi_{\theta}(\cdot|x)} [f(x, y_{1:n})] \\ &= \mathbb{E}_{y_{1:n} \sim \pi_{\theta}(\cdot|x)} \left[\sum_{i=1}^n \nabla_{\theta} \log \pi_{\theta}(y_i | x) (f(x, y_{1:n}) - f(x, y_{1:i-1}, y_{i+1:n})) \right] \end{aligned} \quad (12)$$

[TZSM25] observe that this leads to a gradient estimator that is unbiased. In particular, this makes sense when optimizing objectives such as the pass@ n objective, $f(x, y_{1:n}) = \max_{i \in \{1, \dots, n\}} r(x, y_i)$. In our work on set RL, we choose a different route and instead rely on the marginal set advantage for several reasons.

Firstly, we may wish to optimize objectives that are defined over (or efficiently computable) over a fixed number of n generations. As such, for the general algorithm, we do not make the assumption that the leave-one-out or the leave- k -out baseline is defined. As a straightforward example, the diversity function $d(x, y_{1:n})$ is defined for $n \geq 2$ but does not make sense when $n = 1$. On the other hand, the SPIRAL objective in [HOL+26] is defined over all n but the value of $f_{\text{SPIRAL}}(x, y_{1:n}) = \mathbb{E}_{y_* \sim \pi_{\theta}(\cdot|x, y_{1:n})} [r(x, y_*)]$ can fluctuate in unexpected manners as the number n changes (see [LFC+25] for example) regardless of the quality of individual generations.

Secondly, the leave- k -out estimator implicitly decides on the credit assignment of y_i by attempting to decouple the generations by asking how the value of a set changes as we remove a particular generation from it, whereas marginal set advantage respects the coupling phenomenon in set RL.

We now look at examples that highlight this intuition. For a fixed prompt x , suppress x from the notation and let $\mathcal{G}$ denote the collection of constructed sets of size n . For a generation y_i , write

$$\mathcal{G}(y_i) := \{G \in \mathcal{G} : y_i \in G\}.$$

The marginal set advantage is

$$A_{\text{marg}}(i) := \frac{1}{|\mathcal{G}(y_i)|} \sum_{G \in \mathcal{G}(y_i)} f(G) - \frac{1}{|\mathcal{G}|} \sum_{G \in \mathcal{G}} f(G).$$

By contrast, a natural leave-one-out difference reward is

$$A_{\text{LOO}}(i) := \frac{1}{|\mathcal{G}(y_i)|} \sum_{G \in \mathcal{G}(y_i)} \left(f(G) - f(G \setminus \{y_i\}) \right).$$

This leave-one-out estimator asks whether y_i is pivotal for the value of the sets in which it appears. The marginal set advantage instead asks whether sets containing y_i score above average.

Example 1: pass@2 objective. Let our sampled pool of generations be denoted by

$$\{C_1, C_2, W_1, W_2\},$$

where C denotes a correct generation and W denotes an incorrect generation. Take set size $n = 2$ and define

$$f(G) = \mathbf{1}\{\text{at least one element of } G \text{ is correct}\}.$$

The six possible sets and their scores are

G	$f(G)$
$\{C_1, C_2\}$	1
$\{C_1, W_1\}$	1
$\{C_1, W_2\}$	1
$\{C_2, W_1\}$	1
$\{C_2, W_2\}$	1
$\{W_1, W_2\}$	0

We can calculate the baseline in marginal set advantage as:

$$\bar{f} := \frac{1}{|\mathcal{G}|} \sum_{G \in \mathcal{G}} f(G) = \frac{5}{6}.$$

Now, the marginal set advantage of a correct generation is:

$$A_{\text{marg}}(C) = 1 - \frac{5}{6} = \frac{1}{6} > 0.$$

For an incorrect generation,

$$A_{\text{marg}}(W) = \frac{2}{3} - \frac{5}{6} = -\frac{1}{6} < 0.$$

On the other hand, the leave-one-out advantage estimator for C_1 is computed by considering each G in $\mathcal{G}(C_1)$:

$$\begin{aligned} \{C_1, C_2\} : & 1 - 1 = 0, \\ \{C_1, W_1\} : & 1 - 0 = 1, \\ \{C_1, W_2\} : & 1 - 0 = 1. \end{aligned}$$

Therefore

$$A_{\text{LOO}}(C) = \frac{2}{3} > 0.$$

However, for an incorrect generation W , leaving it out never decreases the pass@2 score. Therefore,

$$A_{\text{LOO}}(W) = 0.$$

Therefore, leave-one-out encourages correct generations with a larger norm advantage but gives no direct negative signal to incorrect generations.

Example 2: majority-vote / self-consistency objective. Let the sampled pool be

$$\{C_1, C_2, C_3, C_4, W_1, W_2, W_3, W_4\},$$

with four correct generations and four incorrect generations. Let $n = 5$, and define a majority-vote objective

$$f(G) = \mathbf{1}\{\# \text{ correct generations in } G \text{ is } \geq 3\}.$$

There are $\binom{8}{5} = 56$ sets. Exactly 28 of them have a correct majority, so

$$\bar{f} = \frac{28}{56} = \frac{1}{2}.$$

For a fixed correct generation C , the number of size-5 sets containing C is $\binom{7}{4} = 35$. Such a set succeeds if the remaining four elements contain at least two additional correct generations. Therefore,

$$\frac{1}{|\mathcal{G}(C)|} \sum_{G \in \mathcal{G}(C)} f(G) = \frac{\binom{3}{2}\binom{4}{2} + \binom{3}{3}\binom{4}{1}}{\binom{7}{4}} = \frac{22}{35}.$$

As such, the marginal set advantage of a correct generation is:

$$A_{\text{marg}}(C) = \frac{22}{35} - \frac{1}{2} = \frac{9}{70} > 0.$$

The marginal set advantage of a fixed incorrect generation W is:

$$\frac{1}{|\mathcal{G}(W)|} \sum_{G \in \mathcal{G}(W)} f(G) = \frac{\binom{4}{3}\binom{3}{1} + \binom{4}{4}\binom{3}{0}}{\binom{7}{4}} = \frac{13}{35}.$$

Therefore

$$A_{\text{marg}}(W) = \frac{13}{35} - \frac{1}{2} = -\frac{9}{70} < 0.$$

Leave-one-out behaves differently. To make the comparison with the marginal set advantage precise, we average the leave-one-out difference over all size-5 sets containing a fixed generation. Since leave-one-out evaluates the objective after removing one element, we extend the majority-vote objective to sets of arbitrary size m by defining

$$f_m(G) := \mathbf{1} \left\{ \# \text{ correct generations in } G > \frac{m}{2} \right\}.$$

Thus, for the original size-5 sets,

$$f_5(G) = \mathbf{1} \{ \# \text{ correct generations in } G \geq 3 \},$$

and for size-4 leave-one-out sets,

$$f_4(G) = \mathbf{1} \{ \# \text{ correct generations in } G \geq 3 \}.$$

For a generation y_i , define the averaged leave-one-out advantage

$$A_{\text{LOO}}(y_i) := \frac{1}{|\mathcal{G}(y_i)|} \sum_{G \in \mathcal{G}(y_i)} (f_5(G) - f_4(G \setminus \{y_i\})).$$

First consider a fixed correct generation C . A set $G \in \mathcal{G}(C)$ contains C and four additional generations chosen from the remaining three correct generations and four incorrect generations. Let t denote the number of additional correct generations among these four. Then

$$t \in \{0, 1, 2, 3\},$$

and the number of such sets with exactly t additional correct generations is

$$\binom{3}{t} \binom{4}{4-t}.$$

The original size-5 set succeeds if

$$1 + t \geq 3,$$

whereas the leave-one-out size-4 set $G \setminus \{C\}$ succeeds if

$$t \geq 3.$$

Therefore

$$f_5(G) - f_4(G \setminus \{C\}) = \mathbf{1} \{1 + t \geq 3\} - \mathbf{1} \{t \geq 3\} = \mathbf{1} \{t = 2\}.$$

Hence only those sets with exactly two additional correct generations contribute nonzero leave-one-out advantage. Therefore

$$\begin{aligned} A_{\text{LOO}}(C) &= \frac{\binom{3}{2} \binom{4}{2}}{\binom{7}{4}} \\ &= \frac{18}{35} > 0. \end{aligned}$$

Now consider a fixed incorrect generation W . A set $G \in \mathcal{G}(W)$ contains W and four additional generations chosen from the remaining four correct generations and three incorrect generations. Let t denote the number of correct generations among these four additional generations. Then

$$t \in \{0, 1, 2, 3, 4\},$$

and the number of such sets is

$$\binom{4}{t} \binom{3}{4-t}.$$

The original size-5 set succeeds if

$$t \geq 3,$$

since the fixed generation W is incorrect. After removing W , the remaining size-4 set also succeeds if

$$t \geq 3.$$

Thus, for every set $G \in \mathcal{G}(W)$,

$$f_5(G) - f_4(G \setminus \{W\}) = \mathbf{1}\{t \geq 3\} - \mathbf{1}\{t \geq 3\} = 0.$$

Therefore

$$A_{\text{LOO}}(W) = 0.$$

Thus the averaged leave-one-out calculation gives

$$A_{\text{LOO}}(C) = \frac{18}{35} > 0, \quad A_{\text{LOO}}(W) = 0.$$

This differs from the marginal set advantage, which gave

$$A_{\text{marg}}(C) = \frac{9}{70} > 0, \quad A_{\text{marg}}(W) = -\frac{9}{70} < 0.$$

Therefore, in this majority-vote example, the marginal set advantage encourages correct generations and discourages incorrect generations, while the raw leave-one-out difference reward encourages correct generations but gives no direct negative signal to incorrect generations. Equivalently, leave-one-out rewards correct generations only when they are pivotal for the majority outcome. In this example, this happens precisely when the size-5 set contains exactly three correct generations, so that removing the fixed correct generation flips the outcome from a correct majority to a non-correct majority.

Example 3: Stylized SPIRAL. The SPIRAL objective is as follows:

$$f_{\text{SPIRAL}}(x, G) = \mathbb{E}_{y_\star \sim \pi_\theta(\cdot | x, G)} [r(x, y_\star)],$$

where G is a set of search traces and $y_\star$ is a downstream aggregation trace conditioned on those search traces. To isolate the credit-assignment effect, consider the following stylized objective. Let

$$C = \text{candidate solution trace}, \quad V = \text{useful lemma}, \quad J = \text{junk trace}.$$

Suppose the aggregator succeeds only when it sees both a candidate trace and at least one verifier trace:

$$f(G) = \mathbf{1}\{C \in G \text{ and some } V \in G\}.$$

Take set size $n = 3$ and sampled pool

$$\{C, V_1, V_2, J\}.$$

The four possible sets are

G	$f(G)$
$\{C, V_1, V_2\}$	1
$\{C, V_1, J\}$	1
$\{C, V_2, J\}$	1
$\{V_1, V_2, J\}$	0

Now, our baseline in marginal set advantage is:

$$\bar{f} = \frac{3}{4}.$$

For the candidate trace,

$$A_{\text{marg}}(C) = 1 - \frac{3}{4} = \frac{1}{4} > 0.$$

For a verifier trace,

$$A_{\text{marg}}(V) = \frac{2}{3} - \frac{3}{4} = -\frac{1}{12} < 0.$$

For the junk trace,

$$A_{\text{marg}}(J) = \frac{2}{3} - \frac{3}{4} = -\frac{1}{12} < 0.$$

Leave-one-out gives a different signal. For a verifier trace V_1 ,

$$\begin{aligned} \{C, V_1, V_2\} : \quad & 1 - 1 = 0, \\ \{C, V_1, J\} : \quad & 1 - 0 = 1, \\ \{V_1, V_2, J\} : \quad & 0 - 0 = 0. \end{aligned}$$

Therefore

$$A_{\text{LOO}}(V_1) = \frac{1}{3} > 0.$$

For the candidate trace,

$$A_{\text{LOO}}(C) = 1 > 0.$$

For the junk trace,

$$A_{\text{LOO}}(J) = 0.$$

Thus leave-one-out gives credit assignment as:

$$C \uparrow, \quad V \uparrow, \quad J \text{ neutral}.$$

There is a notable difference in learning signal here. Marginal set advantage asks which trace types appear in above-average sets; leave-one-out asks which traces are locally indispensable. In this example, verifier traces are sometimes pivotal, so leave-one-out rewards them, even though the marginal set advantage sees that the current pool is already saturated with verifier traces and needs more candidate traces.

Example 4: Polychromic objective. We now give an example using the polychromic objective. For a set $G = \{y_1, \dots, y_n\}$, define

$$f_{\text{poly}}(G) := \bar{r}(G) \cdot d(G),$$

where

$$\bar{r}(G) := \frac{1}{n} \sum_{y \in G} r(y), \quad d(G) := \frac{|\{C(y) : y \in G\}|}{n}.$$

Here $r(y) \in \{0, 1\}$ is the correctness reward and $C(y)$ denotes the reasoning-strategy cluster of generation y . Let the sampled pool be

$$\{C_1, C_2, E, J\},$$

and take set size $n = 3$. Suppose

$$r(C_1) = r(C_2) = 1, \quad r(E) = r(J) = 0.$$

Assume that C_1, C_2 , and J all belong to the same common reasoning cluster, while E belongs to a different exploratory cluster:

$$C(C_1) = C(C_2) = C(J) = a, \quad C(E) = b.$$

Thus E is an incorrect but exploratory generation, while J is an incorrect non-exploratory generation.

There are $\binom{4}{3} = 4$ possible sets:

G	$\bar{r}(G)$	$d(G)$	$f_{\text{poly}}(G)$
$\{C_1, C_2, E\}$	$\frac{2}{3}$	$\frac{2}{3}$	$\frac{4}{9}$
$\{C_1, C_2, J\}$	$\frac{2}{3}$	$\frac{1}{3}$	$\frac{2}{9}$
$\{C_1, E, J\}$	$\frac{1}{3}$	$\frac{2}{3}$	$\frac{2}{9}$
$\{C_2, E, J\}$	$\frac{1}{3}$	$\frac{2}{3}$	$\frac{2}{9}$

Therefore the set-level baseline is

$$\bar{f}_{\text{poly}} = \frac{1}{4} \left(\frac{4}{9} + \frac{2}{9} + \frac{2}{9} + \frac{2}{9} \right) = \frac{5}{18}.$$

We first compute the marginal set advantage. For a correct generation, say C_1 , the sets containing C_1 have scores

$$\frac{4}{9}, \quad \frac{2}{9}, \quad \frac{2}{9}.$$

Thus

$$\begin{aligned} A_{\text{marg}}(C_1) &= \frac{1}{3} \left(\frac{4}{9} + \frac{2}{9} + \frac{2}{9} \right) - \frac{5}{18} \\ &= \frac{8}{27} - \frac{5}{18} \\ &= \frac{1}{54} > 0. \end{aligned}$$

By symmetry,

$$A_{\text{marg}}(C_2) = \frac{1}{54} > 0.$$

For the exploratory incorrect generation E , the sets containing E also have scores

$$\frac{4}{9}, \quad \frac{2}{9}, \quad \frac{2}{9}.$$

Therefore

$$\begin{aligned} A_{\text{marg}}(E) &= \frac{1}{3} \left(\frac{4}{9} + \frac{2}{9} + \frac{2}{9} \right) - \frac{5}{18} \\ &= \frac{8}{27} - \frac{5}{18} \\ &= \frac{1}{54} > 0. \end{aligned}$$

So the marginal set advantage encourages E even though $r(E) = 0$. This happens because E contributes a new reasoning cluster and therefore appears in above-average polychromic sets.

For the non-exploratory incorrect generation J , all sets containing J have score $2/9$. Hence

$$\begin{aligned} A_{\text{marg}}(J) &= \frac{2}{9} - \frac{5}{18} \\ &= -\frac{1}{18} < 0. \end{aligned}$$

Thus the marginal set advantage gives

$$C_1, C_2 \uparrow, \quad E \uparrow, \quad J \downarrow.$$

Now compare this with a raw leave-one-out difference reward. Since removing an element changes the set size, define the size- m version of the polychromic objective by

$$f_{\text{poly},m}(H) := \left(\frac{1}{m} \sum_{y \in H} r(y) \right) \cdot \left(\frac{|\{C(y) : y \in H\}|}{m} \right).$$

For y_i , define the averaged leave-one-out advantage

$$A_{\text{LOO}}(y_i) := \frac{1}{|\mathcal{G}(y_i)|} \sum_{G \in \mathcal{G}(y_i)} (f_{\text{poly},3}(G) - f_{\text{poly},2}(G \setminus \{y_i\})).$$

First consider C_1 . The three leave-one-out terms are

$$\begin{aligned} \{C_1, C_2, E\} : \quad & \frac{4}{9} - f_{\text{poly},2}(\{C_2, E\}) = \frac{4}{9} - \frac{1}{2} = -\frac{1}{18}, \\ \{C_1, C_2, J\} : \quad & \frac{2}{9} - f_{\text{poly},2}(\{C_2, J\}) = \frac{2}{9} - \frac{1}{4} = -\frac{1}{36}, \\ \{C_1, E, J\} : \quad & \frac{2}{9} - f_{\text{poly},2}(\{E, J\}) = \frac{2}{9} - 0 = \frac{2}{9}. \end{aligned}$$

Therefore

$$\begin{aligned} A_{\text{LOO}}(C_1) &= \frac{1}{3} \left(-\frac{1}{18} - \frac{1}{36} + \frac{2}{9} \right) \\ &= \frac{5}{108} > 0. \end{aligned}$$

By symmetry,

$$A_{\text{LOO}}(C_2) = \frac{5}{108} > 0.$$

Now consider the exploratory incorrect generation E . The three leave-one-out terms are

$$\begin{aligned} \{C_1, C_2, E\} : \quad & \frac{4}{9} - f_{\text{poly},2}(\{C_1, C_2\}) = \frac{4}{9} - \frac{1}{2} = -\frac{1}{18}, \\ \{C_1, E, J\} : \quad & \frac{2}{9} - f_{\text{poly},2}(\{C_1, J\}) = \frac{2}{9} - \frac{1}{4} = -\frac{1}{36}, \\ \{C_2, E, J\} : \quad & \frac{2}{9} - f_{\text{poly},2}(\{C_2, J\}) = \frac{2}{9} - \frac{1}{4} = -\frac{1}{36}. \end{aligned}$$

Therefore

$$\begin{aligned} A_{\text{LOO}}(E) &= \frac{1}{3} \left(-\frac{1}{18} - \frac{1}{36} - \frac{1}{36} \right) \\ &= -\frac{1}{27} < 0. \end{aligned}$$

Thus leave-one-out discourages E , even though the marginal set advantage encourages E .

Finally, for the non-exploratory incorrect generation J , we have

$$\begin{aligned} \{C_1, C_2, J\} : \quad & \frac{2}{9} - f_{\text{poly},2}(\{C_1, C_2\}) = \frac{2}{9} - \frac{1}{2} = -\frac{5}{18}, \\ \{C_1, E, J\} : \quad & \frac{2}{9} - f_{\text{poly},2}(\{C_1, E\}) = \frac{2}{9} - \frac{1}{2} = -\frac{5}{18}, \\ \{C_2, E, J\} : \quad & \frac{2}{9} - f_{\text{poly},2}(\{C_2, E\}) = \frac{2}{9} - \frac{1}{2} = -\frac{5}{18}. \end{aligned}$$

Therefore

$$A_{\text{LOO}}(J) = -\frac{5}{18} < 0.$$

In summary, the marginal set advantage gives

$$A_{\text{marg}}(C_1) = A_{\text{marg}}(C_2) = \frac{1}{54}, \quad A_{\text{marg}}(E) = \frac{1}{54}, \quad A_{\text{marg}}(J) = -\frac{1}{18}.$$

Thus

$$\text{marginal:} \quad C_1, C_2 \uparrow, \quad E \uparrow, \quad J \downarrow.$$

The raw leave-one-out difference reward gives

$$A_{\text{LOO}}(C_1) = A_{\text{LOO}}(C_2) = \frac{5}{108}, \quad A_{\text{LOO}}(E) = -\frac{1}{27}, \quad A_{\text{LOO}}(J) = -\frac{5}{18}.$$

Thus

$$\text{leave-one-out:} \quad C_1, C_2 \uparrow, \quad E \downarrow, \quad J \downarrow.$$

This example isolates the optimistic-exploration behavior of the polychromic marginal set advantage. The incorrect exploratory generation E is encouraged because it participates in high-scoring sets by contributing a new reasoning cluster. By contrast, the leave-one-out difference reward asks whether deleting E lowers the value of the current set. Since deleting E also increases the average reward of the remaining set, the leave-one-out signal for E becomes negative. Thus, under the polychromic objective, marginal set advantage and leave-one-out can induce opposite updates for an incorrect but exploratory generation.

3. SCAFFOLD AND MARGINAL ADVANTAGE

3.1. SCAFFOLD VALUE

In this section, we analyse the learning dynamics of set reinforcement learning. In particular, we are interested in understanding the characteristics of generations that are encouraged versus discouraged in the set RL framework. We start by defining the following useful construction:

Definition 3.1 (Scaffold Value). *The scaffold value of a set of generations, $y_{1:n}$, under a policy π and a set-RL objective $f : \mathcal{X} \times \mathcal{Y}^n \rightarrow \mathbb{R}$ is defined to be*

$$\Lambda_f(y_{1:n}; \pi, x) := \text{Cov}_{y'_{1:n} \sim \pi(\cdot|x)}(f(x, y'_{1:n}), \frac{1}{I(a_{1:n})} \sum_{i,j=1}^n 1\{a'_i = a_j\}) \quad (13)$$

where $I(a_{1:n})$ is the maximum size of the intersection of $y_{1:n}$ with any other set $y'_{1:n}$, i.e. $I(a_{1:n}) = \max_{a'_{1:n}} (\sum_{i,j} 1\{a'_i = a_j\})$. We refer to the space $\Lambda_f(\pi) = \{(a_{1:n}, \Lambda_f(a_{1:n}; \pi)) \mid a_{1:n} \in \mathcal{A}^n\}$ as the scaffold of the policy π .

Theorem 3.2 (Scaffold Value of a Homogeneous Set). *Let $f : \mathcal{X} \times \mathcal{Y}^n \rightarrow \mathbb{R}$ be a symmetric set-RL objective. For a homogeneous set $y_{1:n} = y^{\oplus n} = (y, \dots, y)$, the scaffold value under policy π_θ simplifies to:*

$$\Lambda_f(y^{\oplus n}; \pi, x) = \pi_\theta(y \mid x) (\mathbb{E}_{y_{2:n} \sim \pi_\theta(\cdot|x)} [f(x, y, y_{2:n})] - \mathbb{E}_{y_{1:n} \sim \pi_\theta(\cdot|x)} [f(x, y_{1:n})]) \quad (14)$$

Proof. Using the identity $\text{Cov}(A, B) = \mathbb{E}[AB] - \mathbb{E}[A]\mathbb{E}[B]$, we can expand this expression into two terms:

$$\Lambda_f(y^{\oplus n}; \pi, x) = \underbrace{\mathbb{E}_{y'_{1:n}} \left[f(x, y'_{1:n}) \left(\frac{1}{n} \sum_{i=1}^n 1\{y'_i = y\} \right) \right]}_{\text{Term 1}} - \underbrace{\mathbb{E}_{y'_{1:n}} [f(x, y'_{1:n})] \cdot \mathbb{E}_{y'_{1:n}} \left[\frac{1}{n} \sum_{i=1}^n 1\{y'_i = y\} \right]}_{\text{Term 2}} \quad (15)$$

First, we simplify Term 2. By the linearity of expectation and the fact that the samples y'_i are drawn i.i.d., Term 2 becomes:

$$\text{Term 2} = \pi_\theta(y \mid x) \mathbb{E}_{y'_{1:n} \sim \pi_\theta} [f(x, y'_{1:n})] \quad (16)$$

Next, we evaluate Term 1. Applying the linearity of expectation to the sum gives:

$$\text{Term 1} = \frac{1}{n} \sum_{i=1}^n \pi_\theta(y \mid x) \mathbb{E}_{y_{2:n} \sim \pi_\theta} [f(x, y, y_{2:n})] = \pi_\theta(y \mid x) \mathbb{E}_{y_{2:n} \sim \pi_\theta} [f(x, y, y_{2:n})] \quad (17)$$

Finally, recombining Term 1 and Term 2 and factoring out $\pi_\theta(y \mid x)$ provides the desired result. $\square$

Note that $I(a_{1:n})$ is needed to ensure that the second term in the covariance is also bounded between 0 and 1. If a set $a_{1:n}$ is homogeneous, then this is $I(a_{1:n}) = n^2$, and if the set is heterogeneous, then $I(a_{1:n}) = n$.

3.2. MARGINAL SET ADVANTAGE

While this can be useful in analysis, there is a more natural method of understanding the value of a single generation, which we call the marginal set advantage of a generation:

Definition 3.3 (Marginal Set Advantage). *The marginal set advantage of a generation $y \in \mathcal{Y}$ under a policy π and a set-RL objective f is defined as*

$$A_{\text{marg}}^{\#}(x, y; \pi) = \mathbb{E}_{y_{2:n} \sim \pi_{\theta}(\cdot | x)} [f(x, y, y_{2:n})] - \mathbb{E}_{y_{1:n} \sim \pi_{\theta}(\cdot | x)} [f(x, y_{1:n})]. \quad (18)$$

Observation. *Note that $\Lambda_f(y^{\oplus n}; \pi, x) = \pi(y | x) A_{\text{marg}}^{\#}(x, y; \pi)$. For a given generation y , the scaffold value is often a more useful object to study. For example, if $\pi(y | x) = 0$, even if the marginal set advantage of y could be large, the policy would not be learning to upweight that generation since it fails to sample it in the first place. As such, the scaffold value is often a more useful indication of how well a policy might learn an action. We formalize this analysis next.*

The marginal advantage directly determines which generations' log likelihood increases in a set RL algorithm. Suppose, our policy is parametrized as a softmax distribution:

$$\pi_{\theta}(y | x) = \frac{\exp(z_{xy})}{\sum_{y'} \exp(z_{xy'})}$$

where z_{xy} is the output logit of y from x . Then, we have the following result determining how the logit changes for a specific generation.

Lemma 3.4 (Logit shift in set RL). *In set reinforcement learning, assuming the objective function f is symmetric in its arguments and the learning rate is α , the shift in logit after one step gradient update from π_{θ}^k to π_{θ}^{k+1} can be written as:*

$$z_{xy}^{k+1} - z_{xy}^k = \alpha n \pi_{\theta^k}(y | x) A_{\text{marg}}^{\#}(x, y; \pi_{\theta}^k) - \alpha C(\theta^k), \quad (19)$$

where

$$C(\theta^k) = \text{Cov}_{y'_{1:n} \sim \pi_{\theta^k}(\cdot | x)} \left(f(x, y'_{1:n}), \sum_{i=1}^n \pi_{\theta^k}^k(y'_i | x) \right).$$

There is an analogous version of this statement that is concerned with the probability of sampling a set of actions instead:

Proposition 3.5. *Consider any set of actions $y_{1:n} \in \mathcal{Y}^n$. The change in the log probability of sampling this set of actions after one policy update using set RL can be written as the following first-order approximation:*

$$\log \pi_{\theta}^{k+1}(y_{1:n} | x) \approx \log \pi_{\theta}^k(y_{1:n} | x) + \lambda \Lambda_f(y_{1:n}; \pi_{\theta}^k) - \lambda C(\theta^k)$$

where λ is a constant, and

$$C(\theta^k) = \text{Cov}_{y'_{1:n} \sim \pi_{\theta^k}(\cdot | x)} \left(f(x, y'_{1:n}), \sum_{i=1}^n \pi_{\theta^k}^k(y'_i | x) \right).$$

We now compute the marginal set advantages under some objectives to better understand their learning dynamics.

Example 3.6. Let $f(x, y_{1:n}) = \max_{i \in \{1, \dots, n\}} r(x, y_i)$. Suppose that the probability of sampling a successful generation is p . Then, one can easily show that for $f(x, y_{1:n}) = \max_{i \in \{1, \dots, n\}} r(x, y_i)$ and for a fixed incorrect generation y (i.e. $r(x, y) = 0$),

$$\begin{aligned} A_{\text{marg}}^{\#}(x, y) &= \mathbb{E}_{y_{2:n} \sim \pi_{\theta}(\cdot | x)} [f(x, y, y_{2:n})] - \mathbb{E}_{y_{1:n} \sim \pi_{\theta}(\cdot | x)} [f(x, y_{1:n})] \\ &= (1 - (1 - p)^{n-1}) - (1 - (1 - p)^n) \\ &= (1 - p)^n - (1 - p)^{n-1}, \end{aligned}$$

which is non-positive for any $p \in [0, 1]$. As such, under this objective, incorrect generations' log likelihood is never increased by our set RL algorithm.

Example 3.7. We now consider majority voting where the objective is defined by $f(x, y_{1:n}) = r(x, \text{maj}(y_{1:n}))$. Once again, one can show that the marginal set advantage of any incorrect generation is not only always less than or equal to that of a correct generation, it is actually always non-positive.

Example 3.8. Let $f(x, y_{1:n}) = 1\{\exists y_i \neq y_j \mid y_i, y_j \in y_{1:n}\}$. In other words, f is a function that measures pairwise similarity and returns 1 if it finds at least one pair that is not similar. Then, under a uniform policy, $\Lambda_f(y^{\oplus n}; x, \pi)$ is 0 for all y . However, for any non-uniform policy, this is not 0. In particular, for a non-uniform policy, the marginal set advantage is larger for actions that have smaller probability of being sampled from the current distribution. However, the scaffold value multiplies the marginal advantage with the probability of being sampled which counteracts.

3.3. ENTROPY ANALYSIS

Furthermore, the notion of scaffold values helps us understand how entropy changes during set RL gradient updates as well.

Proposition 3.9. Consider the set RL setup on $x \in \mathcal{X}$. After one update to the policy, the change in entropy, $\Delta = \mathcal{H}(\pi_{\theta}^{k+1} | x) - \mathcal{H}(\pi_{\theta}^k | x)$, is given by

$$\Delta \approx -\alpha \text{Cov}_{y_{1:n}} \left(\frac{1}{n} \sum_{i=1}^n \log \pi_{\theta}^k(y_i | x), \text{Cov}_{y'_{1:n}} \left(f(x, y'_{1:n}), \sum_{i,j=1}^n 1\{y_i = y'_j\} \right) \right),$$

where both covariances are taken with respect to $\pi_{\theta}^k(\cdot | x)$ and α is the learning rate.

Proof. First, we show the following change in logits:

$$\begin{aligned} z_{sa}^{k+1} - z_{sa}^k &= \alpha \frac{\partial}{\partial z_{sa}^k} \mathbb{E}_{a_{1:n} \sim \pi_{\theta}^k(\cdot | s)} [f(s, a_{1:n})] \\ &= \alpha \mathbb{E}_{a_{1:n} \sim \pi_{\theta}^k(\cdot | s)} \left[\frac{\partial}{\partial z_{sa}^k} \log(\mathbb{P}(a_{1:n} | s)) f(s, a_{1:n}) \right] \\ &= \alpha \mathbb{E}_{a_{1:n} \sim \pi_{\theta}^k(\cdot | s)} \left[\sum_{i=1}^n \frac{\partial}{\partial z_{sa}^k} \log(\pi_{\theta}^k(a_i | s)) f(s, a_{1:n}) \right] \\ &= \alpha \mathbb{E}_{a_{1:n} \sim \pi_{\theta}^k(\cdot | s)} \left[\sum_{i=1}^n (1\{a_i = a\} - \pi_{\theta}^k(a | s)) f(s, a_{1:n}) \right] \end{aligned}$$

Then, following [CZC⁺25], we have the following first order approximation

$$\begin{aligned} \Delta &\approx \alpha n \mathbb{E}[f(s, a_{1:n})] \text{Cov}_{a \sim \pi_{\theta}^k(\cdot | s)} (\log \pi_{\theta}^k(a | s), \pi_{\theta}^k(a | s)) \\ &\quad - \alpha \text{Cov}_{a \sim \pi_{\theta}^k(\cdot | s)} \left(\log \pi_{\theta}^k(a | s), \mathbb{E}_{a_{1:n} \sim \pi_{\theta}^k(\cdot | s)} \left[f(s, a_{1:n}) \sum_{i=1}^n 1\{a_i = a\} \right] \right) \end{aligned}$$

Note that

$$\begin{aligned} \text{Cov}_{a_{1:n}} \left(f(s, a_{1:n}), \sum_{i=1}^n 1\{a_i = a\} \right) &= \mathbb{E}_{a_{1:n} \sim \pi_{\theta}^k(\cdot | s)} \left[f(s, a_{1:n}) \sum_{i=1}^n 1\{a_i = a\} \right] \\ &\quad - n\pi_{\theta}^k(a | s) \mathbb{E}_{a_{1:n} \sim \pi_{\theta}^k(\cdot | s)} [f(s, a_{1:n})] \end{aligned}$$

Using this, we get that

$$\begin{aligned} \Delta &\approx \alpha n \mathbb{E}[f(s, a_{1:n})] \text{Cov}_{a \sim \pi_{\theta}^k(\cdot | s)} (\log \pi_{\theta}^k(a | s), \pi_{\theta}^k(a | s)) \\ &\quad - \alpha \text{Cov}_{a \sim \pi_{\theta}^k(\cdot | s)} \left(\log \pi_{\theta}^k(a | s), \mathbb{E}_{a_{1:n} \sim \pi_{\theta}^k(\cdot | s)} \left[f(s, a_{1:n}) \sum_{i=1}^n 1\{a_i = a\} \right] \right) \\ &= -\alpha \text{Cov}_{a \sim \pi_{\theta}^k(\cdot | s)} \left(\log \pi_{\theta}^k(a | s), \text{Cov}_{a_{1:n}} \left(f(s, a_{1:n}), \sum_{i=1}^n 1\{a_i = a\} \right) \right). \end{aligned}$$

All that remains to show is

$$\begin{aligned} \text{Cov}_{a_{1:n}} \left(\frac{1}{n} \sum_{i=1}^n \log \pi_{\theta}^k(a_i | s), \text{Cov}_{a'_{1:n}} \left(f(s, a'_{1:n}), \sum_{i,j=1}^n 1\{a'_i = a'_j\} \right) \right) \\ = \text{Cov}_{a \sim \pi_{\theta}^k(\cdot | s)} \left(\log \pi_{\theta}^k(a | s), \text{Cov}_{a_{1:n}} \left(f(s, a_{1:n}), \sum_{i=1}^n 1\{a_i = a\} \right) \right). \end{aligned}$$

This can be seen using the linearity of covariance and the fact that each a_i in $a_{1:n}$ is sampled independently:

$$\begin{aligned} \text{Cov}_{a_{1:n}} \left(\frac{1}{n} \sum_{i=1}^n \log \pi_{\theta}^k(a_i | s), \text{Cov}_{a'_{1:n}} \left(f(s, a'_{1:n}), \sum_{i,j=1}^n 1\{a'_i = a'_j\} \right) \right) \\ = \frac{1}{n} \sum_{i=1}^n \text{Cov}_{a_{1:n}} \left(\log \pi_{\theta}^k(a_i | s), \text{Cov}_{a'_{1:n}} \left(f(s, a'_{1:n}), \sum_{j=1}^n 1\{a_i = a'_j\} \right) \right) \\ = \text{Cov}_{a \sim \pi_{\theta}^k(\cdot | s)} \left(\log \pi_{\theta}^k(a | s), \text{Cov}_{a_{1:n}} \left(f(s, a_{1:n}), \sum_{i=1}^n 1\{a_i = a\} \right) \right). \end{aligned}$$

□

3.4. ANALYSIS OF POLYCHROMIC OBJECTIVES

In this section, we analyse the polychromic objective in detail. Recall that the polychromic objective is defined as

$$f_{\text{poly}}(x, y_1, \dots, y_n) = \frac{1}{n} \sum_{i=1}^n r(x, y_i) d(x, y_1, \dots, y_n)$$

where $d(x, y_1, \dots, y_n)$ is a diversity function. First, we analyse the marginal set advantage of a generation (see eq. (18)) under the polychromic objective. Fix a generation $y \in \mathcal{Y}$ and assume that $\pi_{\theta}(y | x) > 0$. Recall that in set reinforcement learning, a generation y receives increased probability mass (i.e., $\log \pi_{\theta}^{k+1}(y | x) - \log \pi_{\theta}^k(y | x) > 0$) if and only if it has a positive marginal set advantage and non-zero probability under the rollout policy. The marginal set advantage of y can be decomposed as follows:

$$\begin{aligned}
A_{\text{marg}}^{\#}(x, y; f_{\text{poly}}, \theta^k) &= \underbrace{\left(\frac{1}{n} r(x, y) + \frac{n-1}{n} \mathbb{E}_Y[r(x, Y)] \right) \mathbb{E}_{Y_{2:n}}[d(x, y, Y_{2:n})]}_{\text{Term 1: Mean reward of sets containing } y \times \text{Mean diversity of sets containing } y} - \mathbb{E}_Y[r(x, Y)] \mathbb{E}_{Y_{1:n}}[d(x, Y_{1:n})] \\
&+ \underbrace{\text{Cov}_{Y_{2:n}}\left(\frac{1}{n} r(x, y) + \frac{1}{n} \sum_{i=2}^n r(x, Y_i), d(x, y, Y_{2:n})\right)}_{\text{Term 2: Exploration-Exploitation synergy in sets containing } y} - \text{Cov}_{Y_{1:n}}(r(x, Y_1), d(x, Y_{1:n})). \quad (20)
\end{aligned}$$

This equation can be simplified further but the current form is quite explicit about the terms we want to interpret:

- *Term 1* reflects the expected contribution of generation y to the product of reward and diversity in sets that contain it. Importantly, this term depends on both the reward and diversity contributions of other generations in the set. Consequently, even if y itself does not contribute reward (or diversity), this term can still be positive provided other sampled generations contribute reward (or diversity) in expectation.
- *Term 2* captures how the presence of generation y affects the covariance between reward and diversity across sets containing y . Intuitively, this term favors generations that enable sets to simultaneously achieve high reward and high diversity. Thus, generations that help align exploration with exploitation receive positive learning signal.

Therefore, even when $r(x, y) = 0$, a generation can still receive increased probability mass if it contributes diversity while other generations contribute reward, or if it helps produce sets in which reward and diversity are positively aligned. This directly reflects the two desiderata that motivate our method. First, the objective supports *optimistic exploration*: trajectories that attempt novel reasoning strategies can receive positive learning signal even before they become individually successful. Second, the objective *intrinsically encodes the exploration-exploitation trade-off*: the advantage assigned to a trajectory depends in part on whether it helps the model form sets that jointly achieve high reward and high diversity, rather than on a manually tuned weighting coefficient. In this sense, the policy update itself internalizes the balance between exploration and exploitation.

This has a few features.

- In set reinforcement learning with polychromic objectives, as a model’s accuracy improves, exploratory generations can still receive positive learning signal even when incorrect. Because the set objective shares credit among all generations in a set, incorrect but diverse generations can benefit from the presence of correct generations in the same set. Consequently, the incentive for exploration does not vanish as success rate improves.
- Furthermore, the algorithm naturally assigns advantage based on how a generation contributes to aligning exploration and exploitation across sets. This behavior arises from two components: the set RL framework and the multiplicative structure of the polychromic objective - both of which are necessary for the covariance term to emerge in the marginal set advantage.

Note that this analysis can be extended to any objective of the form $f(x, y_{1:n}) = \varphi(x, y_{1:n})\psi(x, y_{1:n})$. Insofar as these are valid objectives for set reinforcement learning that do not collapse into standard reinforcement learning, the algorithmic framework will encourage generations that satisfy both (even if they do not strictly satisfy the other, as long as other generations can compensate) while also strongly encouraging generations that help the model achieve a positive covariance between φ and ψ . One can imagine objectives of this form where $\varphi(x, y_{1:n}) = \frac{1}{n} \sum_{i=1}^n r_A(x, y_i)$ and $\psi(x, y_{1:n}) = \frac{1}{n} \sum_{i=1}^n r_B(x, y_i)$, where r_A and r_B are two different reward functions that we want our model to satisfy in tandem.

We now come to the generalized definition of the polychromic objective. The motivation for the polychromic objective comes from understanding the entropy change in set reinforcement learning as shown in Proposition 3.9. We see that entropy collapses when there is a positive covariance between the log probabilities of generations in a set and the scaffold value of the generations. Note that if this covariance is always positive, we would get entropy collapse. The idea of a polychromic objective is to ensure that the entropy is never allowed to collapse on a set of generations that is a singleton. In other words, the policy can never collapse onto sampling a single generation.

Definition 3.10. A polychromic objective is a function

$$\varphi : \mathcal{X} \times \mathcal{Y}^n \rightarrow \mathbb{R}$$

that factors as

$$\varphi(x, y_{1:n}) = \varphi^{(R)}(x, y_{1:n}) \varphi^{(d)}(x, y_{1:n}),$$

where $\varphi^{(R)}$ and $\varphi^{(d)}$ are scalar-valued functions satisfying:

1. $\text{Cov}_{y_{1:n} \sim \pi_\theta(\cdot|x)}(\varphi^{(R)}(x, y_{1:n}), \sum_{i=1}^n r(x, y_i)) > 0$,
2. $\text{Cov}_{y_{1:n} \sim \pi_\theta(\cdot|x)}(\varphi^{(d)}(x, y_{1:n}), \sum_{i=1}^n 1\{y_i = y\}) < 0$ for any y , and
3. $\varphi^{(R)}(x, \cdot)$ and $\varphi^{(d)}(x, \cdot)$ share the same range, i.e.,

$$\inf \varphi^{(R)}(x, \cdot) = \inf \varphi^{(d)}(x, \cdot), \quad \sup \varphi^{(R)}(x, \cdot) = \sup \varphi^{(d)}(x, \cdot),$$

where the infimum and supremum are taken over all sets $y_{1:n}$.

The practical polychromic objective we have seen before aims to ensure the second requirement using a diversity function since it is, often, intractable to find an objective that satisfies the requirement for all $y \in \mathcal{Y}$.

We confirm now, through a preliminary analysis, that polychromic objectives indeed satisfy the motivation behind it.

Proposition 3.11. For any homogeneous set $y_{1:n} = \{y\}$ where $r(x, y) = 1$, there exists $\epsilon \in (0, 1)$ such that $\Lambda_{f_{\text{poly}}}(y^{\oplus n}; \pi, x) < 0$ when $\pi_\theta(y | x) > \epsilon$. Furthermore, the scaffold values of these homogeneous sets satisfy the bound $\Lambda_{f_{\text{poly}}}(y; \pi, x) \leq \sqrt{\frac{p(1-p)}{n}}$.

This result shows that once a successful generation y accumulates sufficient probability mass, the polychromic objective automatically prevents further entropy collapse onto sets that only contain this action and, when we use larger sets in the set RL framework, the upper bound on the scaffold value of a homogeneous set decreases. This is desirable since it suggests that our policy learns to generate this action without incessantly collapsing probability mass on such homogeneous sets.

Proposition 3.12. Suppose $a_{1:n}$ is heterogeneous where each a_i is unique with probability $\frac{p}{n}$ where $p \in (0, \frac{1}{n})$. Suppose exactly q of the n actions satisfy $r(s, a_i) = 1$, and that any other action $a' \notin a_{1:n}$ with $\pi_\theta(a' | s) > 0$ yields $r(s, a') = 0$. Then, the scaffold value of $a_{1:n}$ satisfies $\Lambda_{f_{\text{poly}}}(a_{1:n}) > \frac{qp^n(1-p)}{n} (1 - (1 - \frac{1}{n})^n)$.

Note that the set in this proposition includes unsuccessful generations as well that contribute diversity. As such, there are, likely, several such heterogeneous sets (provided that the size of sets n is large enough) with positive scaffold values that attract more probability mass than homogeneous sets. Furthermore, the lower bound guarantee increases as the number of successful actions in the set increases. The polychromic objective therefore channels entropy collapse toward those sets that balance success and exploration, rather than permitting collapse onto homogeneous behaviours.

We will now generalize some of the tools used in proving properties of the polychromic objective. In particular, in proving 3.11 and 3.12, we observe that there is a particularly interesting way in which the scaffold value of an action depends on the model's ability to sample other actions. This is yet another indication of how set RL couples generations within a set. The construction of heterogeneous sets aims to make this analysis more rigorous.

4. ASCENDANT VALUE AND OVERSAMPLING

For this section, define the scaffold value as follows:

Definition 4.1 (Scaffold Value). *The scaffold value of a set of generations, $y_{1:n}$, under a policy π and a set-RL objective $f : \mathcal{X} \times \mathcal{Y}^n \rightarrow \mathbb{R}$ is defined to be*

$$\Lambda_f(y_{1:n}; \pi, x) := \text{Cov}_{y'_{1:n} \sim \pi(\cdot|x)} \left(f(x, y'_{1:n}), \frac{1}{n} \sum_{i,j=1}^n 1\{y'_i = y'_j\} \right) \quad (21)$$

As we seek to understand how actions are coupled under a shared learning signal, the scaffold value alone is insufficient to help us do so. The reason for this is that the scaffold value allows oversampling of actions within the set, i.e. $\frac{1}{n} \sum_{i,j=1}^n 1\{y'_i = y'_j\}$ increases even if $y'_{1:n}$ contains oversampled actions that exist in $y_{1:n}$. We now aim to decompose the scaffold value into terms that reveal the coupling between actions versus the contribution from oversampling some action. We start with the following preliminary decomposition:

Proposition 4.2. *Consider a fixed action $y \in \mathcal{Y}$. The scaffold value of a homogeneous set $y^{\oplus n}$ relates to the expected marginal scaffold value of sets containing y by:*

$$\Lambda_{f,\pi}(y^{\oplus n}) = n \mathbb{E}_{y'_{2:n} \sim \pi(\cdot|x)} [\Lambda_{f,\pi}(y_{1:n})] - (n-1)C(\pi) \quad (22)$$

where

$$C(\pi) = \text{Cov}_{y''_{1:n} \sim \pi(\cdot|x)} \left(f(x, y''_{1:n}), \sum_{i=1}^n \pi(y''_i | x) \right).$$

Proof. Let $Z = \{y, y'_2, \dots, y'_n\}$ be a set where the first element is fixed to y , and the remaining $n-1$ elements are sampled i.i.d. $y'_{2:n} \sim \pi(\cdot | x)$. By definition, the scaffold value of Z is:

$$\Lambda_f(Z; \pi, x) = \text{Cov}_{y''_{1:n}} \left(f(x, y''_{1:n}), \frac{1}{n} \sum_{i=1}^n \sum_{j=1}^n 1\{y''_i = Z_j\} \right)$$

We take the expectation over the distribution of the sampled elements $Z_{2:n}$. Because covariance is linear in its second argument and $y''_{1:n}$ is independent of $Z_{2:n}$, we can pass the expectation directly inside the covariance:

$$\mathbb{E}_{Z_{2:n}} [\Lambda_f(Z; \pi, x)] = \text{Cov}_{y''_{1:n}} \left(f(x, y''_{1:n}), \frac{1}{n} \mathbb{E}_{Z_{2:n}} \left[\sum_{i=1}^n \sum_{j=1}^n 1\{y''_i = Z_j\} \right] \right).$$

We now evaluate the inner expectation. We split the sum over j into the fixed element $Z_1 = y$ and the remaining sampled elements Z_j :

$$\begin{aligned} \mathbb{E}_{Z_{2:n}} \left[\sum_{i=1}^n \left(1\{y''_i = y\} + \sum_{j=2}^n 1\{y''_i = Z_j\} \right) \right] &= \sum_{i=1}^n 1\{y''_i = y\} + \sum_{i=1}^n \sum_{j=2}^n \mathbb{E}_{Z_j} [1\{y''_i = Z_j\}] \\ &= \sum_{i=1}^n 1\{y''_i = y\} + (n-1) \sum_{i=1}^n \pi(y''_i | x). \end{aligned}$$

Substituting this simplified sum back into the covariance and using linearity to split it into two distinct terms gives:

$$\begin{aligned} \mathbb{E}_{Z_{2:n}} [\Lambda_f(Z; \pi, x)] &= \frac{1}{n} \text{Cov}_{y''_{1:n}} \left(f(x, y''_{1:n}), \sum_{i=1}^n 1\{y''_i = y\} \right) \\ &\quad + \frac{n-1}{n} \text{Cov}_{y''_{1:n}} \left(f(x, y''_{1:n}), \sum_{i=1}^n \pi(y''_i | x) \right) \end{aligned} \quad (23)$$

We recognize the second term contains exactly the definition of $C(\pi)$. For the first term, we evaluate the scaffold value of a homogeneous set $y^{\oplus n}$ under our definition:

$$\begin{aligned}\Lambda_f(y^{\oplus n}; \pi, x) &= \text{Cov}_{y''_{1:n}} \left(f(x, y''_{1:n}), \frac{1}{n} \sum_{i=1}^n \sum_{j=1}^n 1\{y''_i = y''_j\} \right) \\ &= \text{Cov}_{y''_{1:n}} \left(f(x, y''_{1:n}), \frac{1}{n} \sum_{i=1}^n n 1\{y''_i = y''_i\} \right) \\ &= \text{Cov}_{y''_{1:n}} \left(f(x, y''_{1:n}), \sum_{i=1}^n 1\{y''_i = y''_i\} \right)\end{aligned}$$

This perfectly matches the covariance in the first term of Equation 23. Substituting $\Lambda_f(y^{\oplus n}; \pi, x)$ and $C(\pi)$ yields:

$$\mathbb{E}_{y'_{2:n}} [\Lambda_f(\{y, y'_{2:n}\}; \pi, x)] = \frac{1}{n} \Lambda_f(y^{\oplus n}; \pi, x) + \frac{n-1}{n} C(\pi)$$

Multiplying the entire equation by n and rearranging to isolate $\Lambda_f(y^{\oplus n}; \pi, x)$ concludes the proof:

$$\Lambda_f(y^{\oplus n}; \pi, x) = n \mathbb{E}_{y'_{2:n} \sim \pi(\cdot | x)} [\Lambda_f(\{y, y'_{2:n}\}; \pi, x)] - (n-1)C(\pi)$$

□

Proposition 4.3. *Let $y_{1:n}$ be a fixed target multiset of size n , and let $y'_{1:n}$ be a multiset of the same size sampled by a policy $\pi(\cdot | x)$. Let $c_{y_{1:n}}(v)$ and $c_{y'_{1:n}}(v)$ denote the frequency of an action $v \in \mathcal{Y}$ in the target and sampled sets, respectively.*

The scaffold value of $y_{1:n}$ decomposes exactly as follows:

$$\Lambda_{f,\pi}(y_{1:n}) = h_{f,\pi}(y_{1:n}) + \Omega_{f,\pi}(y_{1:n}) \quad (24)$$

where

$$h_{f,\pi}(y_{1:n}) := \text{Cov}_{y'_{1:n} \sim \pi} \left(f(x, y'_{1:n}), \frac{1}{n} \sum_{v \in \mathcal{Y}} c_{y_{1:n}}(v) \min(c_{y_{1:n}}(v), c_{y'_{1:n}}(v)) \right) \quad (25)$$

and

$$\Omega_{f,\pi}(y_{1:n}) := \text{Cov}_{y'_{1:n} \sim \pi} \left(f(x, y'_{1:n}), \frac{1}{n} \sum_{v \in \mathcal{Y}} c_{y_{1:n}}(v) \max(0, c_{y'_{1:n}}(v) - c_{y_{1:n}}(v)) \right) \quad (26)$$

Proof. Start by noting that:

$$\sum_{i=1}^n \sum_{j=1}^n 1\{y'_i = y'_j\} = \sum_{v \in \mathcal{Y}} c_{y_{1:n}}(v) c_{y'_{1:n}}(v)$$

For any action v , the number of times it is generated by the policy can be partitioned strictly into the elements that overlap with the target set and the excess elements that oversample it:

$$c_{y'_{1:n}}(v) = \min(c_{y_{1:n}}(v), c_{y'_{1:n}}(v)) + \max(0, c_{y'_{1:n}}(v) - c_{y_{1:n}}(v))$$

Substituting this partition into the scaffold value definition yields:

$$\Lambda_{f,\pi}(y_{1:n}) = \text{Cov}_{y'_{1:n}} \left(f(x, y'_{1:n}), \frac{1}{n} \sum_{v \in \mathcal{Y}} c_{y_{1:n}}(v) \left[\min(c_{y_{1:n}}(v), c_{y'_{1:n}}(v)) + \max(0, c_{y'_{1:n}}(v) - c_{y_{1:n}}(v)) \right] \right)$$

By the linearity of covariance, we can separate this into two distinct terms, passing the $c_{y_{1:n}}(v)$ coefficient directly into both parts:

$$\begin{aligned}\Lambda_{f,\pi}(y_{1:n}) &= \text{Cov}_{y'_{1:n}} \left(f(x, y'_{1:n}), \frac{1}{n} \sum_{v \in \mathcal{Y}} c_{y_{1:n}}(v) \min(c_{y_{1:n}}(v), c_{y'_{1:n}}(v)) \right) \\ &\quad + \text{Cov}_{y'_{1:n}} \left(f(x, y'_{1:n}), \frac{1}{n} \sum_{v \in \mathcal{Y}} c_{y_{1:n}}(v) \max(0, c_{y'_{1:n}}(v) - c_{y_{1:n}}(v)) \right).\end{aligned}$$

□

The scaffold value of a single action has the following decomposition.

$$\begin{aligned}\Lambda_f(y^{\oplus n}) &= n\pi(y \mid x)^{n-1} h_{f,\pi}(y^{\oplus n}) + n\mathbb{E}_{y'_{2:n}} [\Omega_{f,\pi}(y \mid y'_{2:n})] \\ &\quad + n\mathbb{E}_{y'_{2:n}} [h_{f,\pi}(y, y'_{2:n}) \mid y'_{2:n} \neq y^{\oplus n-1}] + n\mathbb{E}_{y'_{2:n}} [\Omega_{f,\pi}(\setminus y \mid y'_{2:n})] \\ &\quad - (n-1)C(\pi)\end{aligned}$$

Here, $y_{1:n}$ is a fixed target multiset of size n , and $y'_{1:n}$ is a multiset of the same size sampled by a policy $\pi(\cdot \mid x)$. Let $c_{y_{1:n}}(v)$ and $c_{y'_{1:n}}(v)$ denote the frequency of an action $v \in \mathcal{Y}$ in the target and sampled sets, respectively. Let $c_{\max} = \max_v c_{y_{1:n}}(v)$ be the maximum frequency of any element in the target set.

Here, the terms' definitions are as follows:

$$h_{f,\pi}(y_{1:n}) := \text{Cov}_{y'_{1:n} \sim \pi} \left(f(x, y'_{1:n}), \frac{1}{n} \sum_{v \in \mathcal{Y}} c_{y_{1:n}}(v) \min(c_{y_{1:n}}(v), c_{y'_{1:n}}(v)) \right) \quad (27)$$

$$\Omega_{f,\pi}(y \mid y'_{2:n}) := \text{Cov}_{y'_{1:n} \sim \pi} \left(f(x, y'_{1:n}), \frac{1}{n} c_{y_{1:n}}(y) \max(0, c_{y'_{1:n}}(y) - c_{y_{1:n}}(y)) \right) \quad (28)$$

$$\Omega_{f,\pi}(\setminus y \mid y'_{2:n}) := \text{Cov}_{y'_{1:n} \sim \pi} \left(f(x, y'_{1:n}), \frac{1}{n} \sum_{v \in \mathcal{Y}, v \neq y} c_{y_{1:n}}(v) \max(0, c_{y'_{1:n}}(v) - c_{y_{1:n}}(v)) \right). \quad (29)$$

Observation. The scaffold value can be written as the following decomposition:

$$\Lambda_{f,\pi}(y_{1:n}) = \text{Asc}_{f,\pi}(y_{1:n}) + \mathcal{D}_{f,\pi}(y_{1:n}),$$

where

$$\text{Asc}_{f,\pi}(y_1 = y, y_{2:n}) := \text{Cov}_{y'_{1:n}} \left(f(x, y'_{1:n}), \frac{1}{n} \sum_{v \in \mathcal{Y}} c_{y_{1:n}}(v) \min(c_{y_{1:n}}(v), c_{y'_{1:n}}(v)) \right),$$

and

$$\mathcal{D}_{f,\pi}(y_1 = y, y_{2:n}) := \text{Cov}_{y'_{1:n}} \left(f(x, y'_{1:n}), \frac{1}{n} \sum_{v \in \mathcal{Y}} c_{y_{1:n}}(v) \max(0, c_{y'_{1:n}}(v) - c_{y_{1:n}}(v)) \right).$$

Here, $c_{y_{1:n}}(v)$ and $c_{y'_{1:n}}(v)$ denote the frequency of an action $v \in \mathcal{Y}$ in the target and sampled sets, respectively. Note that $c_{y_{1:n}}(v), c_{y'_{1:n}}(v) \in \{0, 1, \dots, n\}$.

We can now write the scaffold value of the homogeneous set of y as follows:

$$\Lambda_{f,\pi}(y^{\oplus n}) = n \sum_{y_{2:n}} \pi(y_{2:n}) (\text{Asc}_{f,\pi}(y, y_{2:n}) + \mathcal{D}_{f,\pi}(y, y_{2:n})) - (n-1)C(\pi)$$

where $C(\pi) = \text{Cov}_{\tilde{y}_{1:n}}(f(\tilde{y}_{1:n}), \sum_{i=1}^n \pi(\tilde{y}_i))$.

Proposition 4.4. Fix $y \in \mathcal{Y}$ and let $p_{y'} = \pi(y' \mid x)$. For $a \in \{0, 1\}$ define

$$\Gamma_{n,p}^{(a)}(m) := \mathbb{E}_{B \sim \text{Bin}(n-1,p)} [(a+B)(m-a-B)_+].$$

Then

$$\begin{aligned} \Lambda_{f,\pi}(y^{\oplus n}) &= n \sum_{y_{2:n}} \pi(y_{2:n} \mid x) \text{Asc}_{f,\pi}(y, y_{2:n}) \\ &\quad + \text{Cov}_{\tilde{y}_{1:n} \sim \pi} \left(f(x, \tilde{y}_{1:n}), \Gamma_{n,\pi(y|x)}^{(1)}(c_{\tilde{y}_{1:n}}(y)) \right) \\ &\quad + \sum_{y' \in \mathcal{Y}, y' \neq y} \text{Cov}_{\tilde{y}_{1:n} \sim \pi} \left(f(x, \tilde{y}_{1:n}), \Gamma_{n,\pi(y'|x)}^{(0)}(c_{\tilde{y}_{1:n}}(y')) \right) \\ &\quad - (n-1)C(\pi). \end{aligned}$$

Here, we use the notation that $(z)_+ = \max(0, z)$.

Proof. Let $Y_{2:n} = (Y_2, \dots, Y_n)$ denote an auxiliary tuple sampled i.i.d. from $\pi(\cdot \mid x)$, and define the random target multiset

$$Z := (y, Y_{2:n}).$$

Also let $\tilde{Y}_{1:n} \sim \pi(\cdot \mid x)^{\otimes n}$ denote the sampled set appearing inside the scaffold covariance. Throughout the proof, Z and $\tilde{Y}_{1:n}$ are independent.

By the preliminary homogeneous-scaffold identity,

$$\Lambda_{f,\pi}(y^{\oplus n}) = n \mathbb{E}_{Y_{2:n} \sim \pi} [\Lambda_{f,\pi}(y, Y_{2:n})] - (n-1)C(\pi).$$

Using the decomposition

$$\Lambda_{f,\pi}(y, Y_{2:n}) = \text{Asc}_{f,\pi}(y, Y_{2:n}) + \mathcal{D}_{f,\pi}(y, Y_{2:n}),$$

we obtain

$$\begin{aligned} \Lambda_{f,\pi}(y^{\oplus n}) &= n \mathbb{E}_{Y_{2:n} \sim \pi} [\text{Asc}_{f,\pi}(y, Y_{2:n})] \\ &\quad + n \mathbb{E}_{Y_{2:n} \sim \pi} [\mathcal{D}_{f,\pi}(y, Y_{2:n})] - (n-1)C(\pi). \end{aligned}$$

Since

$$\mathbb{E}_{Y_{2:n} \sim \pi} [\text{Asc}_{f,\pi}(y, Y_{2:n})] = \sum_{y_{2:n}} \pi(y_{2:n} \mid x) \text{Asc}_{f,\pi}(y, y_{2:n}),$$

it remains to rewrite the averaged descendant term.

For a fixed auxiliary target set $Z = (y, Y_{2:n})$, the descendant value is

$$\mathcal{D}_{f,\pi}(Z) = \text{Cov}_{\tilde{Y}_{1:n} \sim \pi} \left(f(x, \tilde{Y}_{1:n}), \frac{1}{n} \sum_{v \in \mathcal{Y}} c_Z(v) (c_{\tilde{Y}_{1:n}}(v) - c_Z(v))_+ \right).$$

Multiplying by n and averaging over Z gives

$$n \mathbb{E}_Z [\mathcal{D}_{f,\pi}(Z)] = \mathbb{E}_Z \left[\text{Cov}_{\tilde{Y}_{1:n} \sim \pi} \left(f(x, \tilde{Y}_{1:n}), \sum_{v \in \mathcal{Y}} c_Z(v) (c_{\tilde{Y}_{1:n}}(v) - c_Z(v))_+ \right) \right].$$

Because Z is independent of $\tilde{Y}_{1:n}$, and covariance is linear in its second argument, the expectation over Z can be passed inside the covariance:

$$\begin{aligned} n \mathbb{E}_Z [\mathcal{D}_{f,\pi}(Z)] &= \text{Cov}_{\tilde{Y}_{1:n} \sim \pi} \left(f(x, \tilde{Y}_{1:n}), \mathbb{E}_Z \left[\sum_{v \in \mathcal{Y}} c_Z(v) (c_{\tilde{Y}_{1:n}}(v) - c_Z(v))_+ \right] \right) \\ &= \sum_{v \in \mathcal{Y}} \text{Cov}_{\tilde{Y}_{1:n} \sim \pi} \left(f(x, \tilde{Y}_{1:n}), \mathbb{E}_Z [c_Z(v) (c_{\tilde{Y}_{1:n}}(v) - c_Z(v))_+] \right). \end{aligned}$$

Now fix an action $v \in \mathcal{Y}$. Since $Z = (y, Y_2, \dots, Y_n)$, we have $c_Z(v) = \mathbf{1}\{v = y\} + \sum_{j=2}^n \mathbf{1}\{Y_j = v\}$. Let $a_v := \mathbf{1}\{v = y\}$ and $p_v := \pi(v \mid x)$. Then $\sum_{j=2}^n \mathbf{1}\{Y_j = v\} \sim \text{Bin}(n-1, p_v)$, and hence $c_Z(v) \stackrel{d}{=} a_v + B_v$ and $B_v \sim \text{Bin}(n-1, p_v)$. Therefore, conditional on the realized sampled set $\tilde{Y}_{1:n}$, if we write $m_v := c_{\tilde{Y}_{1:n}}(v)$, then

$$\begin{aligned} \mathbb{E}_Z \left[c_Z(v) (c_{\tilde{Y}_{1:n}}(v) - c_Z(v))_+ \right] &= \mathbb{E}_{B_v \sim \text{Bin}(n-1, p_v)} \left[(a_v + B_v)(m_v - a_v - B_v)_+ \right] \\ &= \Gamma_{n, p_v}^{(a_v)}(m_v). \end{aligned}$$

Substituting $p_v = \pi(v \mid x)$, $a_v = \mathbf{1}\{v = y\}$, and $m_v = c_{\tilde{Y}_{1:n}}(v)$ yields

$$n \mathbb{E}_Z [\mathcal{D}_{f, \pi}(Z)] = \sum_{v \in \mathcal{Y}} \text{Cov}_{\tilde{Y}_{1:n} \sim \pi} \left(f(x, \tilde{Y}_{1:n}), \Gamma_{n, \pi(v|x)}^{(\mathbf{1}\{v=y\})} (c_{\tilde{Y}_{1:n}}(v)) \right).$$

Combining this identity with the earlier decomposition gives

$$\begin{aligned} \Lambda_{f, \pi}(y^{\oplus n}) &= n \sum_{y_{2:n}} \pi(y_{2:n} \mid x) \text{Asc}_{f, \pi}(y, y_{2:n}) \\ &\quad + \sum_{v \in \mathcal{Y}} \text{Cov}_{\tilde{Y}_{1:n} \sim \pi} \left(f(x, \tilde{Y}_{1:n}), \Gamma_{n, \pi(v|x)}^{(\mathbf{1}\{v=y\})} (c_{\tilde{Y}_{1:n}}(v)) \right) \\ &\quad - (n-1)C(\pi). \end{aligned}$$

Renaming the summation variable v as y' allows us to conclude. $\square$

Definition 4.5 (Ascendant Value, Oversampling Value). *Consider a fixed $y \in \mathcal{Y}$ and a set RL objective f . Then, the ascendant value of y is defined as*

$$\text{Asc}_{f, \pi}(y_1 = y, y_{2:n}) := \text{Cov}_{y'_{1:n}} \left(f(x, y'_{1:n}), \frac{1}{n} \sum_{v \in \mathcal{Y}} c_{y_{1:n}}(v) \min(c_{y_{1:n}}(v), c_{y'_{1:n}}(v)) \right).$$

On the other hand, the oversampling value of y is defined as

$$\text{Ovs}_{f, \pi}(y) = \text{Cov}_{y'_{1:n} \sim \pi} (f(x, y'_{1:n}), \Gamma_{n, \pi(y|x)}^{(1)}(c_{y'_{1:n}}(y)) - \Gamma_{n, \pi(y|x)}^{(0)}(c_{y'_{1:n}}(y)))$$

where

$$\Gamma_{n, p}^{(a)}(m) := \mathbb{E}_{B \sim \text{Bin}(n-1, p)} [(a + B)(m - a - B)_+].$$

Now, we summarize our work so far in the following proposition:

Proposition 4.6. *Consider an action $y \in \mathcal{Y}$ and a set RL objective function f . Then,*

$$\Lambda_{f, \pi}(y^{\oplus n}) = n \mathbb{E}_{y_{2:n} \sim \pi(\cdot|x)} [\text{Asc}_{f, \pi}(y, y_{2:n})] + \text{Ovs}_{f, \pi}(y) + \tilde{C}(\pi) - (n-1)C(\pi), \quad (30)$$

where

$$\tilde{C}(\pi) = \sum_{y' \in \mathcal{Y}} \text{Cov}_{\tilde{y}_{1:n} \sim \pi} \left(f(x, \tilde{y}_{1:n}), \Gamma_{n, \pi(y'|x)}^{(0)}(c_{\tilde{y}_{1:n}}(y')) \right)$$

and

$$C(\pi) = \text{Cov}_{\tilde{y}_{1:n} \sim \pi} (f(x, \tilde{y}_{1:n}), \sum_{i=1}^n \pi(\tilde{y}_i \mid x)).$$

This decomposition of the scaffold value of y (and, therefore, also the marginal set advantage of y) allows us to study its contribution from coupling with other actions and the contribution from oversampling y alone.

5. STRATIFICATION

We now introduce a small set of tools for studying coupling in set RL without requiring knowledge of the policy. The basic idea is to study the geometry of the set objective itself. In particular, we ask: how many elements of a high-value set must be fixed before the objective's attainable range begins to shrink?

Throughout this section, we fix a prompt x and study the objective $f(x, Y_{1:n})$ as a function of the multiset of generations. We write $\mathcal{Y}^{\oplus n}$ for multisets of size n . For two multisets A and B , we write

$$A \preceq B$$

if A is a sub-multiset of B , and we write $A \oplus B$ for multiset union. For a partial multiset $A \in \mathcal{Y}^{\oplus a}$ with $a \leq n$, define the residual attainable range of f after fixing A by

$$\mathcal{R}_f^x(A) := \left\{ f(x, A \oplus Z) : Z \in \mathcal{Y}^{\oplus(n-a)} \right\}.$$

We also define the associated tight interval

$$I_f^x(A) := [L_f^x(A), U_f^x(A)] := [\inf \mathcal{R}_f^x(A), \sup \mathcal{R}_f^x(A)].$$

Therefore, $I_f^x(A)$ is the tight attainable range of the objective after the elements in A have been fixed.

Definition 5.1 (Minimal maximizing subset). *Fix a prompt x , and let*

$$M_f(x) := U_f^x(\emptyset)$$

be the maximum attainable value of f at x . A multiset $T \in \mathcal{Y}^{\oplus m}$ is called a maximizing subset for f at x if

$$I_f^x(T) = \{M_f(x)\}.$$

Equivalently, once the elements of T have been fixed, every completion of the remaining $n - m$ generations attains the maximum value of f .

A maximizing subset T is called minimal if no maximizing subset of strictly smaller cardinality exists. We denote the collection of minimal maximizing subsets by

$$\mathfrak{T}_x(f).$$

When the prompt x is clear from context, we write simply $\mathfrak{T}(f)$.

Definition 5.2 (Degree of stratification). *Let $T \in \mathfrak{T}_x(f)$ be a minimal maximizing subset. The degree of stratification of f along T is*

$$\deg_x(f; T) := \max \left\{ |A| : A \prec T \text{ and } I_f^x(A) = I_f^x(\emptyset) \right\},$$

where $A \prec T$ means that A is a proper sub-multiset of T .

Thus $\deg_x(f; T)$ is the largest number of elements from a minimal maximizing subset that can be revealed while the attainable range of the objective remains unchanged. If even fixing one element of T shrinks the attainable range, then

$$\deg_x(f; T) = 0.$$

Definition 5.3 (Stratification). *A stratification of f at x is a pair*

$$\mathfrak{S}(f(x)) = (T, \deg_x(f; T)), \quad T \in \mathfrak{T}_x(f).$$

When there are multiple minimal maximizing subsets, the degree may in principle depend on the choice of T . In such cases, we define the degree by taking the maximum over minimal subsets:

$$\deg_x(f) := \max_{T \in \mathfrak{T}_x(f)} \deg_x(f; T)$$

and we use the corresponding T in the definition of the stratification.

The degree captures a simple notion of delayed information. If $\deg_x(f; T) = s$, then there exists a proper sub-multiset $A \prec T$ of size s such that fixing A does not shrink the attainable range of the objective. Only after adding more elements from the subset does the objective begin to become more constrained.

Example 5.4 (Pass@ k objective). Let

$$f_{\text{pass}@k}(x, G) := \mathbf{1}\{\exists y \in G \text{ such that } r(x, y) = 1\}.$$

Assume that both successful and unsuccessful generations are attainable. Then

$$I_{f_{\text{pass}@k}}^x(\emptyset) = [0, 1].$$

Any single correct generation y with $r(x, y) = 1$ is a minimal maximizing subset, since

$$I_{f_{\text{pass}@k}}^x(\{y\}) = \{1\}.$$

Therefore

$$T = (y), \quad |T| = 1, \quad \deg_x(f_{\text{pass}@k}; T) = 0.$$

Therefore, pass@ k has singleton stratification: one correct generation is enough to force the maximum, which confirms the intuition that the objective does not couple generations.

Example 5.5 (Standard RL objective written as a set objective). Let

$$f_{\text{RL}}(x, y_{1:n}) := \frac{1}{n} \sum_{i=1}^n r(x, y_i),$$

where $r(x, y) \in [0, 1]$ and the maximum reward 1 is attainable. The maximum value of f_{RL} is 1, and this value is forced only when all n generations have reward 1. Therefore a minimal maximizing subset is

$$T = (y_1, \dots, y_n), \quad r(x, y_i) = 1 \quad \text{for all } i.$$

If $A \preceq T$ has cardinality a , then

$$I_{f_{\text{RL}}}^x(A) = \left[\frac{a}{n}, 1 \right].$$

In particular, fixing even one reward-1 generation changes the attainable range from $[0, 1]$ to $[1/n, 1]$. Hence

$$|T| = n, \quad \deg_x(f_{\text{RL}}; T) = 0.$$

This reflects the fact that the average-reward objective decomposes immediately across individual generations.

Example 5.6 (Majority-voting objective). Let

$$f_{\text{maj}}(x, G) := \mathbf{1}\{\# \text{ correct generations in } G \geq t\},$$

where t is the number of correct generations required for the majority-voted answer to be correct. For strict majority voting,

$$t = \left\lfloor \frac{n}{2} \right\rfloor + 1.$$

If ties are counted as correct when n is even, then one may instead take $t = n/2$.

A minimal maximizing subset is

$$T = (y_1, \dots, y_t), \quad r(x, y_i) = 1 \quad \text{for all } i.$$

Indeed, once t correct generations are fixed, every completion has a correct majority. Therefore

$$I_{f_{\text{maj}}}^x(T) = \{1\}.$$

On the other hand, if $A \prec T$ has cardinality $a < t$, then both success and failure are still attainable: the remaining slots can be completed with enough correct generations to reach the threshold, or with incorrect generations to miss the threshold. Thus

$$I_{f_{\text{maj}}}^x(A) = [0, 1] = I_{f_{\text{maj}}}^x(\emptyset) \quad \text{for all } A \prec T.$$

Consequently,

$$|T| = t, \quad \deg_x(f_{\text{maj}}; T) = t - 1.$$

In the even- n tie-accepting convention, this gives

$$\deg_x(f_{\text{maj}}; T) = \frac{n}{2} - 1.$$

Majority voting therefore has delayed stratification: no strict subset of the minimal majority subset changes the attainable range.

Example 5.7 (Polychromic objective). Consider a polychromic set objective of the form

$$f_{\text{poly}}(x, G) := \bar{r}(x, G) d(x, G),$$

where

$$\bar{r}(x, G) := \frac{1}{n} \sum_{y \in G} r(x, y)$$

is the average reward and $d(x, G) \in [0, 1]$ is a diversity score. For concreteness, suppose that each generation belongs to a reasoning cluster $C(y)$, and take

$$d(x, G) := \frac{1}{n(n-1)} \sum_{\substack{y, y' \in G \\ y \neq y'}} \mathbf{1}\{C(y) \neq C(y')\}.$$

Thus $d(x, G) = 1$ when all generations in the set belong to distinct reasoning clusters, while $d(x, G) = 0$ when all generations collapse to the same cluster.

Assume that there exist n correct generations

$$y_1, \dots, y_n$$

such that

$$r(x, y_i) = 1 \quad \text{for all } i,$$

and

$$C(y_i) \neq C(y_j) \quad \text{for all } i \neq j.$$

Then

$$T = (y_1, \dots, y_n)$$

is a minimal maximizing subset for f_{poly} , since maximum polychromic value requires both perfect average reward and full diversity.

We now compute the degree. If a single element y_i of T is fixed, then the remaining generations can be completed in two extreme ways. They can all collapse to the same cluster as y_i , in which case $d(x, G) = 0$ and $f_{\text{poly}}(x, G) = 0$. Or they can be completed by correct generations in distinct clusters, in which case $f_{\text{poly}}(x, G) = 1$. Hence

$$I_{f_{\text{poly}}}^x(\{y_i\}) = [0, 1] = I_{f_{\text{poly}}}^x(\emptyset).$$

However, if two distinct elements $y_i, y_j \in T$ are fixed, then both average reward and diversity are bounded away from zero for every completion. Thus

$$I_{f_{\text{poly}}}^x(\{y_i, y_j\}) \subsetneq [0, 1].$$

Therefore

$$|T| = n, \quad \deg_x(f_{\text{poly}}; T) = 1.$$

The polychromic objective is therefore more coupled than the average-reward objective: one correct trace alone does not change the attainable range, but two correct and mutually diverse traces already force a nontrivial amount of reward-diversity structure.

Definition 5.8 (Compatible set objectives). *Two objectives f and g are compatible at prompt x if there exist minimal maximizing subsets*

$$T_f \in \mathfrak{T}_x(f), \quad T_g \in \mathfrak{T}_x(g),$$

such that either

$$T_f \preceq T_g \quad \text{or} \quad T_g \preceq T_f.$$

The smaller subset is called the compatible minimal subset, and the larger subset is called the compatible product subset.

For product objectives, it is useful to define a relative degree. The reason is that if $T_f \preceq T_g$, then T_g is also a maximizing subset for f , but it need not be a minimal one. The degree of f measured along T_g can therefore differ from the intrinsic degree of f measured along T_f .

Definition 5.9 (Relative degree along a subset). *Let T be any maximizing subset for f at x , not necessarily a minimal one. The relative degree of f along T is*

$$\deg_x(f; T) := \max \{ |A| : A \prec T \text{ and } I_f^x(A) = I_f^x(\emptyset) \}.$$

When T is minimal, this agrees with the degree of stratification defined above.

We can now state a product rule for compatible objectives.

Proposition 5.10 (Product rule for compatible objectives). *Fix a prompt x . Let f and g be non-constant, bounded objectives satisfying*

$$f(x, \cdot) \geq 0, \quad g(x, \cdot) \geq 0,$$

and suppose

$$I_f^x(\emptyset) = [0, M_f], \quad I_g^x(\emptyset) = [0, M_g],$$

with $M_f, M_g > 0$.

Assume that f and g are compatible. That is, there exist minimal maximizing certificates

$$T_f \in \mathfrak{T}_x(f), \quad T_g \in \mathfrak{T}_x(g),$$

such that, without loss of generality,

$$T_f \preceq T_g.$$

Then T_g is a minimal maximizing subset for the product objective $f \cdot g$. Moreover, the degree of the product along this product subset satisfies

$$\deg_x(f \cdot g; T_g) = \max \{ \deg_x(f; T_g), \deg_x(g; T_g) \},$$

where $\deg_x(f; T_g)$ and $\deg_x(g; T_g)$ denote relative degrees measured along the same subset T_g .

Proof. Let

$$M_f := U_f^x(\emptyset), \quad M_g := U_g^x(\emptyset).$$

Since T_g is a maximizing subset for g , fixing T_g forces

$$g = M_g.$$

Since $T_f \preceq T_g$ and T_f is a maximizing subset for f , fixing T_g also forces

$$f = M_f.$$

Therefore fixing T_g forces

$$f \cdot g = M_f M_g.$$

Because both functions are nonnegative and bounded above by their respective maxima, $M_f M_g$ is the maximum attainable value of the product. Thus T_g is a maximizing subset for $f \cdot g$.

We now show minimality. Suppose A is any maximizing subset for $f \cdot g$. Then for every completion of A ,

$$f(x, \cdot) g(x, \cdot) = M_f M_g.$$

Since

$$0 \leq f(x, \cdot) \leq M_f, \quad 0 \leq g(x, \cdot) \leq M_g,$$

this equality can hold only if, for every completion of A ,

$$f(x, \cdot) = M_f \quad \text{and} \quad g(x, \cdot) = M_g.$$

Hence A is a maximizing subset for both f and g . Therefore

$$|A| \geq |T_f| \quad \text{and} \quad |A| \geq |T_g|.$$

Since $|T_g| \geq |T_f|$, we have

$$|A| \geq |T_g|.$$

Thus no subset for $f \cdot g$ can be shorter than T_g , so T_g is a minimal maximizing subset for the product.

It remains to prove the degree identity. Let $A \prec T_g$. Since A can be completed by the remaining elements of T_g , and T_g contains both a maximizing subset for f and a maximizing subset for g , we have

$$U_f^x(A) = M_f, \quad U_g^x(A) = M_g, \quad U_{f \cdot g}^x(A) = M_f M_g.$$

Thus the only way the tight interval can shrink is through the lower endpoint.

Because f and g are nonnegative,

$$L_{f \cdot g}^x(A) = 0$$

if and only if

$$L_f^x(A) = 0 \quad \text{or} \quad L_g^x(A) = 0.$$

Indeed, if one factor can be made arbitrarily close to zero, then the product can be made arbitrarily close to zero. Conversely, if both lower bounds are strictly positive, then the product is bounded below by their product.

Since

$$I_f^x(\emptyset) = [0, M_f], \quad I_g^x(\emptyset) = [0, M_g], \quad I_{f \cdot g}^x(\emptyset) = [0, M_f M_g],$$

we have

$$I_{f \cdot g}^x(A) = I_{f \cdot g}^x(\emptyset)$$

if and only if

$$I_f^x(A) = I_f^x(\emptyset) \quad \text{or} \quad I_g^x(A) = I_g^x(\emptyset).$$

Therefore the largest proper sub-multiset of T_g along which the product range remains unchanged is the larger of the corresponding largest range-preserving sub-multisets for f and g . Hence

$$\deg_x(f \cdot g; T_g) = \max \{ \deg_x(f; T_g), \deg_x(g; T_g) \}.$$

□

The relative-degree qualification in lemma 5.10 is important. In general, it is not correct to compare the degree of $f \cdot g$ with the intrinsic degrees of f and g computed on different minimal subsets. If $T_f \preceq T_g$, then elements of $T_g \setminus T_f$ may be irrelevant to the minimal stratification of f , but they can still affect the product's range-preservation behavior. The clean product rule holds when all degrees are measured along the same compatible product subset.

6. APPLICATION: EXPLORATION

In this section, we present some examples of algorithms in set reinforcement learning.

6.1. POLY-PPO

Polychromic Proximal Policy Optimisation (POLY-PPO) is an algorithm that aims to do a mix of set reinforcement learning (as in the idealized formulation presented in §2.1) and standard reinforcement learning. Recall that in the idealised formulation, we sample a set of n actions at every state visited.

1. If the initial state of the environment is fixed, we can sample a set of trajectories from the initial state and optimize the policy at the initial state using set reinforcement learning since we will naturally get n actions sampled at the initial state.
2. If we can reset our environment to any desired state, then we can select a set of states from the trajectories we have sampled so far and generate $n - 1$ more trajectories from those states to be able to do set reinforcement learning there. This is called *vine sampling* and was studied by [SLM⁺17] and [KAP⁺25] amongst others.

The algorithm POLY-PPO operates under the assumption that we can reset our environment to any state. At a high-level, the algorithm first samples trajectories using vine sampling. Then, at any state where we have sampled a set (larger than a singleton) of trajectories, the algorithm optimises the policy using set reinforcement learning. At all other states, i.e. states from which we did not sample a set of trajectories, the algorithm optimises the policy using standard reinforcement learning. The pseudocode is provided in algorithm 2. We refer the reader to [HOX⁺26] for implementation details. Note that:

1. The algorithm induces a balance between exploration and exploitation in two different ways. Firstly, at all the rollout states where we can use set reinforcement learning, we encourage a balance between exploration and exploitation by optimising the polychromic objective. Secondly, at all other states, when we apply just the GAE advantage, we are encouraging pure exploitation since the advantage estimator is in the standard RL framework using the standard return maximization objective. For long horizon tasks, since the number of rollout states is much smaller than the non-rollout state, we use a window hyperparameter T such that, at each rollout state s_t , we apply the set RL advantage to all states and actions from s_t up to s_{t+T} .
2. The mix of set reinforcement learning and standard reinforcement learning allows the model to explore-then-exploit *within* each generation. For classical RL environments that do not involve foundation models, this was a useful inductive bias. When training large language models with the polychromic objective, this inductive bias is not required as the model, by itself, learns to explore and exploit both across and within generations. Regardless, vine sampling is a powerful technique that allows us to get closer to the sequential set RL formulation we discussed in §2 where we sample a set of actions at every state visited.

This algorithm was proposed in [HOX⁺26] and the original motivation was to optimise the polychromic objective, however the recipe can be adapted for optimising other set RL objectives as well.

Algorithm 2 Polychromic PPO

```

1: for iteration = 1, 2, ... do
2:   Collect trajectories under  $\pi_\beta$ ; rollout vines  $\tau_{1:N}$  from rollout states
3:   if  $s_t$  is a rollout state then
4:     Construct sets,  $G_1, \dots, G_K$  of  $n$  trajectories from  $s_t$ 
5:     // Given  $s$ , let  $\mathcal{G}(a)_{1:p}$  be all sets containing a trajectory starting with  $a$ 
6:     Set  $\hat{A}(s_t, a_t) = \frac{1}{p} \sum_{i=1}^p f_{\text{poly}}(s_t, \mathcal{G}(a_t)_i) - \frac{1}{K} \sum_{j=1}^K f_{\text{poly}}(s_t, G_j)$ 
7:   else
8:     Compute  $\hat{A}(s_t, a_t)$  via GAE
9:   end if
10:  Update  $\pi_\theta$  for  $K$  epochs on minibatches  $\mathcal{B}$  by maximizing the PPO objective
11:  Set  $\pi_\beta \leftarrow \pi_\theta$ 
12: end for
```

6.2. POLY-EPO

Polychromic Exploratory Policy Optimisation (POLY-EPO) is an algorithm that leverages the general recipe in algorithm 1 for optimizing the practical polychromic objective proposed in [OHR⁺26].

We first discuss the implementation of the diversity function. There are a few considerations to keep in mind. Firstly, we require the computation of the diversity of each set, i.e. $d(x, y_1, \dots, y_n)$, to be computationally practical. This is because our set RL recipe constructs all $K = \binom{N}{n}$ sets of size n from the N generations per prompt; if scoring each set under the polychromic objective is computationally expensive, then the runtime complexity of our algorithm will increase rapidly as we scale up the algorithm to a setting where we sample more generations per prompt. Secondly, we must be able to do this computation for a wide variety of problem settings. Language models are post-trained on problems in mathematics, coding, instruction-following, etc., and the training set for each of these settings can be quite large. We require a method for measuring diversity of sets of generations that can be easily transferred to a wide variety of settings.

The main ingredient we use is a language model judge (LM-judge) that is used to cluster the responses sampled from our policy. Since LM-judges are already incorporated in post-training methods in a wide variety of settings such as proof-writing and non-verifiable domains, this allows POLY-EPO to benefit from the post-training infrastructure that is already widely adopted. Furthermore, clustering responses based on semantic similarity is a substantially simpler task than solving the problem, which allows us to leverage the generator-verifier gap in designing our method.

Recall that, for each prompt x , we first sample N independent generations and, then, construct sets of size n from them. We first take all N generations conditioned on the same prompts and ask an LM-judge to cluster them based on their semantic similarity in their reasoning strategy. Let $\mathcal{C}(y_i) \in \{1, \dots, N\}$ be the cluster to which generation y_i belongs. Then, we define the diversity function for each set $\{y_{(j_1)}, \dots, y_{(j_n)}\}$ to be:

$$d(x, y_{(j_1)}, \dots, y_{(j_n)}) = \frac{\text{No. of clusters in } (y_{(j_1)}, \dots, y_{(j_n)})}{n} = \frac{|\{\mathcal{C}(y_{(j_1)}), \dots, \mathcal{C}(y_{(j_n)})\}|}{n}$$

Note that $d(x, y_{(j_1)}, \dots, y_{(j_n)}) \in [0, 1]$, which is required for the polychromic objective. To increase the reliability of the judge, we use in-context learning to steer the judge to cluster them based on the strategy (both high-level and low-level) used in each generation, as opposed to differences in tone, style, etc. Now, we have all the ingredients for computing the score of each set under the polychromic objective. Since scoring each set is computationally straightforward, we can construct all $K = \binom{N}{n}$ sets in our set RL recipe from algorithm 1.

6.3. SPIRAL

In this section, we discuss SPIRAL [HOL⁺26], a framework for optimizing a model to effectively search and aggregate. In standard reinforcement learning for reasoning, we optimize a model’s usage of sequential compute only: the model samples a chain-of-thought which is rewarded in isolation. SPIRAL allows the model to reason using sequential, parallel and aggregative compute which is optimized end-to-end in a unified manner. This is useful for several reasons. Firstly, at test-time, we tend to scale compute all three axes, which means allowing the model to optimize its usage of all three primitives of inference compute bridges the gap between training time optimization goals and test-time deployment. Secondly, allowing the model to learn to use all three primitives should, in theory at least, enable the model to better balance exploration and exploitation, by understanding how each primitive is better suited to some goal (for e.g. the model can learn that it should use parallel compute to sample exploratory trajectories which can be later aggregated into an optimal output).

The problem itself is formulated as follows:

$$\max_{\theta, \phi} \mathbb{E}_{y_{1:n} \sim \pi_{\theta}(\cdot|x)} \left[\mathbb{E}_{y_* \sim \pi_{\phi}(\cdot|x, y_{1:n})} [r(x, y_*)] \right]. \quad (31)$$




Figure 2: SPIRAL trains a language model to use sequential, parallel, and aggregative inference compute end-to-end. Given an input x , SPIRAL samples parallel reasoning traces, which the LM synthesizes into a final trace y_* . During optimization, SPIRAL uses only the final reward $r(x, y_*)$ to optimize the parallel generations via set RL and the aggregation trace via standard RL. Consequently, the model jointly learns to generate traces that aggregate usefully and to synthesize them effectively.

Optimizing this formulation requires the model to learn how to allocate sequential compute within a single chain of thought (i.e., each $y_i \sim \pi_\theta(\cdot | x)$), how to sample high-quality parallel traces that broadly explore the solution space (i.e., $y_{1:n} \sim \pi_\theta(\cdot | x)$), and how to aggregate these traces effectively to arrive at a correct answer (i.e., $y_* \sim \pi_\phi(\cdot | x)$). This formulation naturally permits using a different model for aggregation ($\phi \neq \theta$) or the same model for both stages ($\phi = \theta$). Throughout the rest of this section, we will refer to the generations $y_{1:n} \sim \pi_\theta(\cdot | x)$ as search traces and the generation $y_* \sim \pi_\phi(\cdot | x, y_{1:n})$ as an aggregation trace.

We now introduce a scalable recipe for optimizing this objective. We start by considering the case where $\pi_\phi = \pi_\theta$ (the algorithm is similar for the case where this is not true). Define the following set-level objective:

$$f_{\text{spiral}}(x, y_{1:n}) = \mathbb{E}_{y_* \sim \pi_\theta(\cdot | x, y_{1:n})} [r(x, y_*)].$$

With this defined, we now see that taking the gradient of eq. (31), we get:

$$\begin{aligned} & \nabla_\theta \mathbb{E}_{y_{1:n} \sim \pi_\theta(\cdot | x)} [\mathbb{E}_{y_* \sim \pi_\theta(\cdot | x, y_{1:n})} [r(x, y_*)]] \\ &= \underbrace{\mathbb{E}_{y_{1:n} \sim \pi_\theta(\cdot | x)} \left[f_{\text{spiral}}(x, y_{1:n}) \sum_{i=1}^n \nabla_\theta \log \pi_\theta(y_i | x) \right]}_{\text{Set RL gradient}} \\ & \quad + \underbrace{\mathbb{E}_{y_{1:n} \sim \pi_\theta(\cdot | x)} [\mathbb{E}_{y_* \sim \pi_\theta(\cdot | x, y_{1:n})} [r(x, y_*) \nabla_\theta \log \pi_\theta(y_* | x, y_{1:n})]]}_{\text{Standard RL gradient}}. \end{aligned} \quad (32)$$

These two gradient terms emerge naturally because we are jointly optimizing both the search traces within a set and the aggregation trace conditioned on that set. The set RL gradient updates the policy to generate an optimal set $y_{1:n}$ that the policy can effectively aggregate. Simultaneously, the standard RL gradient updates the policy to reliably synthesize a given set of generations into the optimal aggregation solution, y_* . As such, this framework captures a co-evolutionary procedure. Next, we discuss a scalable algorithm for optimizing eq. (32).

On-policy Data Collection. During each data collection step, we first sample a batch of prompts $x \sim \mathcal{D}$. We then sample N_1 search traces conditioned on the prompt: $y_1, \dots, y_{N_1} \sim \pi_\theta(\cdot | x)$. Next, we construct sets of size n , following the set reinforcement learning recipe in [OHR⁺26]. In particular, we uniformly sample K sets without replacement, $G_1, \dots, G_K$, from the collection of all $\binom{N_1}{n}$ unordered sets of size n . In other words, each set $G_i = \{y_{i,1}, \dots, y_{i,n}\}$ contains n unique generations and no two sets in $\{G_1, \dots, G_K\}$ are identical. Finally, we sample aggregation traces: conditioned on each set G_i , we sample $y_1^{G_i}, \dots, y_{N_2}^{G_i} \sim \pi_\theta(\cdot | x, y_{i,1}, \dots, y_{i,n})$. Each aggregation trace is evaluated using the reward function




Figure 3: SPIRAL consists of two levels of generation. First, the model samples N_1 search traces, $y_1, \dots, y_{N_1} \sim \pi_{\theta}(\cdot | x)$, from the problem x . It then uniformly samples K sets, each consisting of n search traces. For each set, it samples N_2 aggregation traces conditioned on both the problem and the set G_i , $y_1^{G_i}, \dots, y_{N_2}^{G_i} \sim \pi_{\theta}(\cdot | x, G_i)$. Rewards are evaluated only at the end of this full inference compute pipeline compute pipeline, using the aggregation trace $r(x, y_j^{G_i})$. For the aggregation traces, SPIRAL uses a standard within-set centered advantage, $A(x, y_j^{G_i}) = r(x, y_j^{G_i}) - \frac{1}{N_2} \sum_{k=1}^{N_2} r(x, y_k^{G_i})$. For the search traces, SPIRAL uses set reinforcement learning: each trace y_i receives a marginal set advantage obtained by averaging the advantages of all sampled sets that contain it ($\mathcal{G}(y_i)$), $A_{\text{marg}}^{\#}(x, y_i) = \frac{1}{|\mathcal{G}(y_i)|} \sum_{G \in \mathcal{G}(y_i)} A^{\#}(x, G)$. Thus, search traces are optimized according to how much they help induce successful downstream aggregation traces when used as part of a conditioning set.

$r(x, y_j^{G_i})$ for trace index $j \in \{1, \dots, N_2\}$ and set index $i \in \{1, \dots, K\}$. This data collection pipeline is illustrated in fig. 3.

Policy Updates. The policy gradient of SPIRAL relies exclusively on the reward assigned to the aggregation traces. Each search trace is optimized via set reinforcement learning, and each aggregation trace is optimized via standard reinforcement learning.

To optimize the search traces, we apply the general set reinforcement learning recipe proposed by [OHR⁺26]. First, we score each set G_i under the empirical objective function:

$$\hat{f}_{\text{spiral}}(x, G_i) = \frac{1}{N_2} \sum_{j=1}^{N_2} r(x, y_j^{G_i}), \quad \forall i \in \{1, \dots, K\}.$$

We construct a Monte Carlo estimate of the baseline score across sets:

$$\hat{f}_{\text{spiral}}(x) = \frac{1}{K} \sum_{i=1}^K \hat{f}_{\text{spiral}}(x, G_i).$$

This allows us to compute the set advantage of each set G_i as:

$$A^{\#}(x, G_i; \hat{f}_{\text{spiral}}) = \hat{f}_{\text{spiral}}(x, G_i) - \hat{f}_{\text{spiral}}(x).$$

Finally, we compute the marginal set advantage for each individual search trace:

$$A_{\text{marg}}^{\#}(x, y; f_{\text{spiral}}) = \frac{1}{|\mathcal{G}(y)|} \sum_{G \in \mathcal{G}(y)} A^{\#}(x, G; f_{\text{spiral}}), \quad (33)$$

where $\mathcal{G}(y) = \{G \in \{G_1, \dots, G_K\} \mid y \in G\}$ is the collection of sets containing y .

To optimize the aggregation traces, we compute the standard RL advantage. The advantage of each aggregation trace is computed within its respective set, rather than across all sets:

$$A(x, y_j^{G_i}) = r(x, y_j^{G_i}) - \frac{1}{N_2} \sum_{k=1}^{N_2} r(x, y_k^{G_i}). \quad (34)$$

We observe that using a per-set baseline leads to significantly lower variance, which helps the model learn how to reliably aggregate traces.

Having computed the advantages for each generation across both levels, we can now plug the advantages into standard reinforcement learning algorithms to optimize the model, following the recipe proposed in [OHR⁺26]. In our implementations, we substitute the standard advantage function in REINFORCE with importance sampling.

SPIRAL is another example of a set reinforcement learning objective that couples generations within a set. Insofar as the aggregation trace is a synthesis of the reasoning explored across various search traces, each generation can couple with another. As a relatively naive illustration, one can imagine that search trace y_1 consist of an idea A that does not succeed by itself whereas search trace y_2 consists of another idea B that does not succeed by itself but y_* uses both ideas A and B to solve the problem.

REFERENCES

- [CZC⁺25] Ganqu Cui, Yuchen Zhang, Jiacheng Chen, Lifan Yuan, Zhi Wang, Yuxin Zuo, Haozhan Li, Yuchen Fan, Huayu Chen, Weize Chen, Zhiyuan Liu, Hao Peng, Lei Bai, Wanli Ouyang, Yu Cheng, Bowen Zhou, and Ning Ding. The entropy mechanism of reinforcement learning for reasoning language models, 2025.
- [HOL⁺26] Jubayer Ibn Hamid, Ifdita Hasan Orney, Michael Y Li, Omar Shaikh, Yoonho Lee, Dorsa Sadigh, Chelsea Finn, and Noah Goodman. Spiral: Learning to search and aggregate. *arXiv preprint arXiv:2606.23595*, 2026.
- [HOX⁺26] Jubayer Ibn Hamid, Ifdita Hasan Orney, Ellen Xu, Chelsea Finn, and Dorsa Sadigh. Polychromic objectives for reinforcement learning, 2026.
- [KAP⁺25] Amirhossein Kazemnejad, Milad Aghajohari, Eva Portelance, Alessandro Sordoni, Siva Reddy, Aaron Courville, and Nicolas Le Roux. Vineppo: Refining credit assignment in rl training of llms, 2025.
- [KL02a] Sham Kakade and John Langford. Approximately optimal approximate reinforcement learning. In *Proceedings of the Nineteenth International Conference on Machine Learning, ICML '02*, page 267–274, San Francisco, CA, USA, 2002. Morgan Kaufmann Publishers Inc.
- [KL02b] Sham M. Kakade and John Langford. Approximately optimal approximate reinforcement learning. In *International Conference on Machine Learning*, 2002.
- [LFC⁺25] Zichong Li, Xinyu Feng, Yuheng Cai, Zixuan Zhang, Tianyi Liu, Chen Liang, Weizhu Chen, Haoyu Wang, and Tuo Zhao. Llms can generate a better answer by aggregating their own responses, 2025.
- [LZY⁺25] Tianjian Li, Yiming Zhang, Ping Yu, Swarnadeep Saha, Daniel Khashabi, Jason Weston, Jack Lanchantin, and Tianlu Wang. Jointly reinforcing diversity and quality in language model generations, 2025.
- [OHR⁺26] Ifdita Hasan Orney, Jubayer Ibn Hamid, Shreya S Ramanujam, Shirley Wu, Hengyuan Hu, Noah Goodman, Dorsa Sadigh, and Chelsea Finn. Poly-epo: Training exploratory reasoning models, 2026.
- [SLM⁺17] John Schulman, Sergey Levine, Philipp Moritz, Michael I. Jordan, and Pieter Abbeel. Trust region policy optimization, 2017.
- [SWD⁺17] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. Proximal policy optimization algorithms, 2017.
- [SWZ⁺24] Zhihong Shao, Peiyi Wang, Qihao Zhu, Runxin Xu, Junxiao Song, Xiao Bi, Haowei Zhang, Mingchuan Zhang, Y. K. Li, Y. Wu, and Daya Guo. Deepseekmath: Pushing the limits of mathematical reasoning in open language models, 2024.
- [TZSM25] Yunhao Tang, Kunhao Zheng, Gabriel Synnaeve, and Rémi Munos. Optimizing language models for inference time objectives using reinforcement learning, 2025.
- [WK25] Christian Walder and Deep Karkhanis. Pass@k policy optimization: Solving harder reinforcement learning problems, 2025.